

Lecture 12: Semantic Integrity Constraints

CS 1555: Database Management Systems

Constantinos Costa

<http://db.cs.pitt.edu/courses/cs1555/current.term/>

March 07, 2019, 16:00-17:15
University of Pittsburgh, Pittsburgh, PA



Lectures based: P. Chrysanthis & N. Farnan Lectures

Structural Constraints

o Constraints (on Attributes):

- NOT NULL
- DEFAULT value
 - without the DEFAULT-clause, the default value is NULL
- PRIMARY KEY (attribute-list)
- UNIQUE (attribute list)
 - allows the specification of alternative key
- FOREIGN KEY (key) REFERENCES table (key)



CS 1555: Database Management Systems - Constantinos Costa

Referential Triggered Actions

- Actions if a Referential Integrity constraint is violated
 - SET NULL
 - CASCADE (propagate action)
 - SET DEFAULT
- Qualify actions by the triggering condition:
 - ON DELETE
 - ON UPDATE
- Note: Oracle does not support ON UPDATE & SET DEFAULT



CS 1555: Database Management Systems - Constantinos Costa

Create Table with RI Trigger Actions

CREATE TABLE LIBRARIAN

```
(  Name    name_dom,  
   SSN     ssn_dom,  
   Section INTEGER,  
   Address address_dom,  
   Gender  gender_dom,  
   Birthday DATE,  
   Salary  DEC(8,2),  
  
   CONSTRAINT librarian_PK PRIMARY KEY (SSN) DEFERRABLE,  
   CONSTRAINT librarian_FK  
   FOREIGN KEY (Section) REFERENCES SECTION (SNO)  
   On Delete SET DEFAULT On Update CASCADE  
   DEFERRABLE );
```



CS 1555: Database Management Systems - Constantinos Costa

Semantic Integrity Constraints

- A constraint is expressed as a Predicate, a condition similar to the one at the WHERE-clause of a query
- Three DDL constructs
 - Checks
 - Assertions
 - Triggers



CS 1555: Database Management Systems - Constantinos Costa

Check constraints

```
CREATE TABLE DEPARTMENT (  
    Dname                VARCHAR(15) UNIQUE,  
    Dnumber              INT PRIMARY KEY,  
    Created_date         DATE,  
    Budget               DECIMAL(15,2),  
    Mgr_ssn              CHAR(9) NOT NULL,  
    Mgr_start_date       DATE,  
    CONSTRAINT dept_fk   FOREIGN KEY (Mgr_ssn) REFERENCES EMPLOYEE(Ssn)  
    CONSTRAINT dept_budget_IC1  
        CHECK ((Budget >= 0) AND (Budget IS NOT NULL)),  
    CONSTRAINT mgr_IC  
        CHECK (Created_date <= Mgr_start_date)  
);
```



CS 1555: Database Management Systems - Constantinos Costa

6

Assertions

- **CREATE OR REPLACE ASSERTION** *aname*
CHECK *predicate*;
- **DROP ASSERTION** *aname*;
- **CREATE OR REPLACE ASSERTION** *budget_constraint*
CHECK (**NOT EXISTS**
 (**SELECT** * **FROM** DEPARTMENT D **WHERE** Budget
 <(**SELECT** SUM(Salary) **FROM** EMPLOYEE E
 WHERE E.Dno = D.Dnumber)))
- Assertion predicates often built around EXISTS and NOT EXISTS
- Note: PostgreSQL does not support assertions



Creating triggers

- **CREATE TRIGGER** *trig_name* *time event*
ON *table_name*
 [**FOR EACH** { **ROW** | **STATEMENT** }]
 [**WHEN** (*condition*)]
 EXECUTE PROCEDURE *func_name* ();
- Triggers can be dropped with:
 - **DROP TRIGGER** [**IF EXISTS**] *trig_name*
 ON *table_name* [**CASCADE** | **RESTRICT**];



Trigger example

```
CREATE TRIGGER Name_Trim  
BEFORE INSERT  
ON Student  
FOR EACH ROW  
EXECUTE PROCEDURE trim_name();
```



When triggers can fire

- *time*
 - BEFORE
 - AFTER
 - INSTEAD OF
- *event*
 - INSERT
 - DELETE
 - UPDATE [OF *att_name* [, ...]]
 - TRUNCATE



Compatibility

WHEN	EVENT	ROW	STATEMENT
BEFORE	INSERT, UPDATE, DELETE	Tables	Tables and views
	TRUNCATE	—	Tables
AFTER	INSERT, UPDATE, DELETE	Tables	Tables and views
	TRUNCATE	—	Tables
INSTEAD OF	INSERT, UPDATE, DELETE	Views	—
	TRUNCATE	—	—



A few more trigger points

- FOR EACH ROW classifies the trigger as a *row trigger*
 - In contrast to a *statement trigger*
 - E.g., if an SQL UPDATE updates 10 rows, a row trigger fires 10 times, a statement trigger fires once
- NEW and OLD allow reference to be made to new and old values of tuples affected by the trigger (and satisfy the WHEN clause) in a row trigger
- WHEN specifies a condition that must be met for the trigger to fire



Trigger maintenance

- **ALTER TABLE** *tname*

{ ENABLE | DISABLE } [*trigger_name* | **ALL** | **USER**];



Trigger procedures

- How can we define procedures in SQL??
 - We can't, we need a procedural language to do this
 - PL/pgSQL
 - The procedural language for PostgreSQL



PL/pgSQL general structure

```
CREATE FUNCTION func_name() RETURNS r_type AS $$  
[ DECLARE  
    declarations ]  
BEGIN  
    statements  
END;  
$$ LANGUAGE plpgsql;  
  
DROP FUNCTION [IF EXISTS] func_name() [CASCADE|RESTRICT];
```



If statements

- IF *condition*
 THEN
 statement
 ...
 ELSIF *condition*
 THEN
 statement
 ...
 ELSE
 statement
 ...
END IF;



Iteration

- **LOOP**
statement
...
EXIT; *-- or: EXIT WHEN condition;*
END LOOP;
- **WHILE** *condition* **LOOP**
statement
...
END LOOP;
- **FOR** *counter* **IN** [**REVERSE**] *val1..val2* [**BY** *exp*]
LOOP
statement
...
END LOOP;

Comment

**Integer loop variant,
PostgreSQL supports several
other variants**



Variables

- **Declaration:**
var_name [**CONSTANT**] *type* [**NOT NULL**] [{**DEFAULT** | **:=**} *expression*];
- **Aliasing:**
 - *new_name* **ALIAS FOR** *old_name*
- **Types can be copied from existing tables:**
 - *table_name.row_name*%**TYPE**



Processing relations

- How can we represent relations in a procedural context?

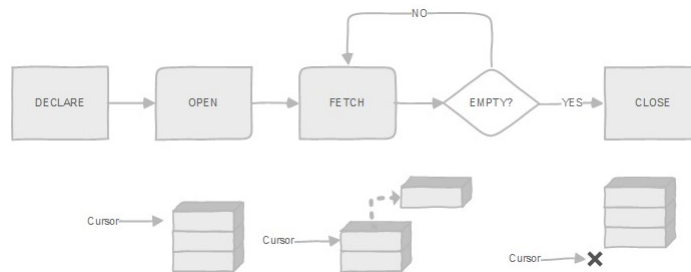


Records: storing table rows

- *rec_name* **RECORD**;
 - Has no predefined structure
 - Substructure is set when it is assigned a value
- *rec_name table_name%ROWTYPE*;
 - Structured to match the schema of *table_name*



Cursors



Source: <http://www.postgresqltutorial.com/plpgsql-cursor/>



CS 1555: Database Management Systems - Constantinos Costa

21 of 16

Cursors

- `cur_name [[NO] SCROLL] CURSOR [(args)] FOR query;`
 - E.g.:
 - `curs1 CURSOR FOR SELECT * FROM table1;`
 - `curs2 CURSOR (key integer) FOR`
`SELECT * FROM table1 WHERE att1 = key;`
- Before a cursor can be used, it must be opened
 - `OPEN curs1;`
 - `OPEN curs2(42);`
 - `OPEN curs2(key := 42);`
- Should be closed when finished:
 - `CLOSE cur_name;`



CS 1555: Database Management Systems - Constantinos Costa

22

FETCHing and MOVEing

- **FETCH** [*direction* { **FROM** | **IN** }] *cursor* **INTO** *target*;
 - *target* should be a **RECORD** or list of variables
 - **RECORD** can be a specific **ROWTYPE**
 - *direction* can take on many forms, e.g.:
 - **FETCH** curs1 **INTO** rowvar;
 - **FETCH** curs2 **INTO** foo, bar, baz;
 - **FETCH** **LAST** **FROM** curs3 **INTO** x, y;
 - **FETCH** **RELATIVE** -2 **FROM** curs4 **INTO** recvar;
 - Special variable **FOUND** will be set to true if a row is returned from the fetch
- **MOVE** [*direction* { **FROM** | **IN** }] *cursor*;
 - **MOVE** *direction* has all the flexibility of **FETCH** *direction*



Cursor example

```
CREATE FUNCTION gpa_summer() RETURNS INTEGER AS $$
DECLARE
    gpa_sum INTEGER := 0;
    st_cursor CURSOR FOR
        SELECT ID, Name, Major, GPA FROM Students;
    student_rec Students%ROWTYPE;
BEGIN
    OPEN st_cursor;
    LOOP
        FETCH st_cursor INTO student_rec;
        IF NOT FOUND THEN
            EXIT;
        END IF;
        gpa_sum := gpa_sum + student_rec.GPA;
    END LOOP;
    CLOSE st_cursor;
    RETURN gpa_sum;
END;
$$ LANGUAGE plpgsql;
```



Exception handling in PL/pgSQL

- Add an EXCEPTION clause at the end of a pgSQL block:

- **BEGIN**

statements

...

EXCEPTION

WHEN *condition* [**OR** *condition* ...] **THEN**
handler_statements

...

[**WHEN** *condition* [**OR** *condition* ...] **THEN**
handler_statements

...

...]

END;



CS 1555: Database Management Systems - Constantinos Costa

25

(Toy) exception example

```
CREATE TABLE db (a INT PRIMARY KEY, b TEXT);
CREATE FUNCTION fun(key INT, data TEXT) RETURNS VOID AS $$
BEGIN
    LOOP
        UPDATE db SET b = data WHERE a = key;
        IF found THEN
            RETURN;
        END IF;
        BEGIN
            INSERT INTO db(a,b) VALUES (key, data);
            RETURN;
        EXCEPTION WHEN unique_violation THEN
            -- Do nothing, but ponder how this could ever
            happen...
        END;
    END LOOP;
END;
$$
LANGUAGE plpgsql;
```



CS 1555: Database Management Systems - Constantinos Costa

26

Getting back to triggers...

- PL/pgSQL code that will run when a trigger is fired must be stored in a *trigger procedure*
 - Has a return type of TRIGGER
 - Takes no arguments
 - Will have access to special variables
 - NEW
 - OLD
 - etc.
 - Can return replacement rows or NULL to indicate that an INSERT or UPDATE operation should not go through



Assertions and Triggers

- **CREATE OR REPLACE ASSERTION** budget_constraint
CHECK (NOT EXISTS
 (SELECT * FROM DEPARTMENT D WHERE
 Budget < **(SELECT SUM(Salary) FROM EMPLOYEE E**
 WHERE E.Dno = D.Dnumber))));
- How can we create a trigger to accomplish the same functionality?



Solution

- **CREATE FUNCTION** blocker() RETURNS TRIGGER AS \$\$
BEGIN
 IF (EXISTS
 (SELECT * FROM DEPARTMENT D WHERE Budget <
 (SELECT SUM(Salary) FROM EMPLOYEE E
 WHERE E.Dno = D.Dnumber)))
 THEN
 RETURN NULL;
 END IF;
 RETURN NEW;
END;
\$\$ LANGUAGE plpgsql;
- **CREATE TRIGGER** test1
 BEFORE INSERT OR UPDATE OF Budget ON DEPARTMENT
 FOR EACH ROW EXECUTE PROCEDURE blocker();



To ensure action is taken in either case

- **CREATE TRIGGER** test2
 BEFORE INSERT OR UPDATE OF Salary ON
 EMPLOYEE
 FOR EACH ROW EXECUTE PROCEDURE blocker();



Statement-level trigger example

- **CREATE FUNCTION** Logger() **RETURNS TRIGGER AS \$\$**
BEGIN
 INSERT INTO Audit_Log
 VALUES ('Students', current_timestamp);
END;
\$\$ LANGUAGE plpgsql;
- **CREATE TRIGGER** Audit_Updater
 AFTER INSERT OR DELETE OR UPDATE ON Students
 EXECUTE PROCEDURE Logger();



CS 1555: Database Management Systems - Constantinos Costa

31

Row-level trigger example

- **CREATE FUNCTION** Add_DL() **RETURNS TRIGGER AS \$\$**
BEGIN
 INSERT INTO DL **VALUES** (NEW.ID, NEW.GPA);
 RETURN NEW;
END;
\$\$ LANGUAGE plpgsql;
- **CREATE TRIGGER** trig_deans_list
 AFTER INSERT ON STUDENTS
 FOR EACH ROW
 WHEN (NEW.GPA > 3.5)
 EXECUTE PROCEDURE Add_DL();



CS 1555: Database Management Systems - Constantinos Costa

32

A bad trigger

- **CREATE FUNCTION** increment() **RETURNS TRIGGER AS \$\$**
BEGIN
 SELECT MAX(ID) + 1 **INTO** NEW.ID **FROM** STUDENTS;
 RETURN NEW;
END;
\$\$ **LANGUAGE** plpgsql;
- **CREATE TRIGGER** bad_auto_sid
 AFTER INSERT ON Students
 FOR EACH ROW
 EXECUTE PROCEDURE increment();



An improved trigger

- **CREATE FUNCTION** increment() **RETURNS TRIGGER AS \$\$**
BEGIN
 SELECT MAX(ID) + 1 **INTO** NEW.ID **FROM** STUDENTS;
 RETURN NEW;
END;
\$\$ **LANGUAGE** plpgsql;
- **CREATE TRIGGER** good_auto_sid
 BEFORE INSERT ON Students
 FOR EACH ROW
 EXECUTE PROCEDURE increment();



In closing

- CHECK constraints and ASSERTIONS follow a declarative approach to integrity constraint enforcement
- TRIGGERS take a procedural approach

