

# Lecture 14: Data Storage

## CS 1555: Database Management Systems

Constantinos Costa

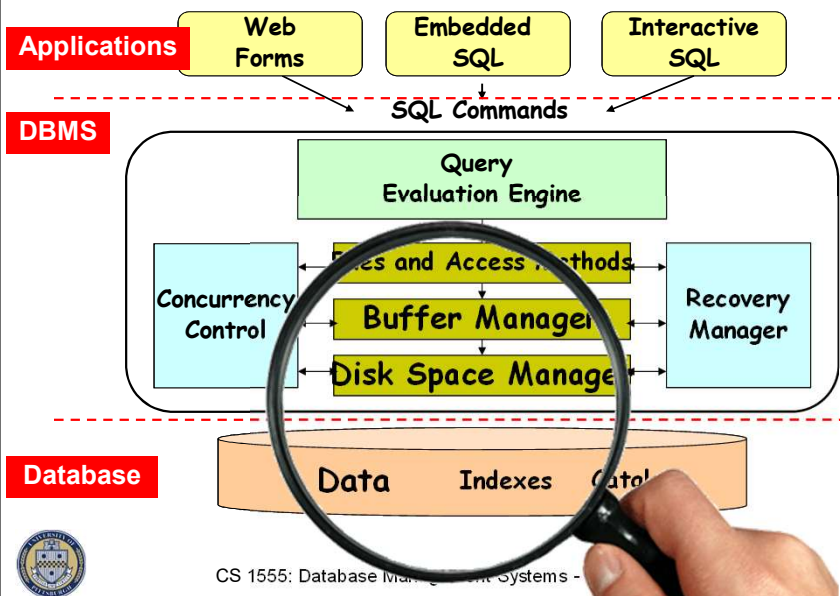
<http://db.cs.pitt.edu/courses/cs1555/current.term/>

March 21, 2019, 16:00-17:15  
University of Pittsburgh, Pittsburgh, PA



Lectures based: P. Chrysanthis & N. Farnan Lectures

## Database Management System (DBMS)



CS 1555: Database Management Systems -

## Data Storage

- Two DBMS fundamental questions?
  1. How do we store and manage very large volumes of data?
  2. What representation and data structures best support efficient manipulation of data?
    - RAM model vs. I/O model of Computation



CS 1555: Database Management Systems - Constantinos Costa

## Storage Hierarchy

- Primary Storage: random access; volatile
  - *Cache* - on board or level-2 cache
  - *Main memory*
- Secondary storage: random access; non-volatile
  - Flash-based or Solid State Disk
  - *Magnetic disk*
    - Virtual Memory (Main-memory DBS), File System, DBMS
- Tertiary storage: non-volatile
  - *Optical disk* / juke boxes – random access
  - *Magnetic cartridge* / tape silos – seq. access

Speed; \$\$



CS 1555: Database Management Systems - Constantinos Costa

## Storage Options

- ❓ Disk is slow (order of **millisecond**),  
but RAM is fast (order of **nanosecond**)
  
- ❓ Why not store everything in RAM?
  1. Expensive cost
  2. Small capacity
  3. Volatile



CS 1555: Database Management Systems - Constantinos Costa

## Disks and Files

- DBMS stores information on hard disks
  
- This has major implications on DBMS design:
  - **READ**: transfer data from disk to RAM
  - **WRITE**: transfer data from RAM to disk
  - Both are high-cost (I/O) operations (i.e., slow)
    - must be planned carefully!



CS 1555: Database Management Systems - Constantinos Costa

## Magnetic Disks

- Data is stored and retrieved in units called disk **blocks** or **pages**
- Main advantage over tapes:  
**random access** (vs. *sequential*)




- Unlike RAM, **time** to retrieve a page **varies** depending upon location on disk



CS 1555: Database Management Systems - Constantinos Costa

## Magnetic Disks

- *Disk block* is the unit of transfer: a disk *block* is a continuous sequence of sectors from a single track of one platter.
- Physical address of  $x$ :  $\langle \text{volume\#}, \text{platter\#}, \text{track\#}, \text{sector\#} \rangle$
- Block address of  $x$ :  $\langle \text{volume\#}, \text{sector/block\#} \rangle$



CS 1555: Database Management Systems - Constantinos Costa

## Accessing a Disk Block

- Time to access (read/write) a disk block:
  1. **seek time:**
    - moving arms to position disk head on track
  2. **rotational delay:**
    - waiting for block to rotate under head
  3. **transfer time:**
    - actually moving data to/from disk surface



CS 1555: Database Management Systems - Constantinos Costa

## Performance Measures of Disks

For any block of data,

Average block access time =

**seek time + rotational delay + transfer time**


- Seek time and rotational delay dominate!
  - o **Disk Access** = *Seek time* + *Rotation latency*  
**<2000** 8 - 20 ms = 2-10 ms + 4-10 ms  
**now** 2 -10 ms = 1-6 ms + 1-4 ms
  - o **Data Transfer Rate** = 4-8MB/sec (max 25-40MB/sec)
    - Block transfer is less than 0.5 msec
    - Multiple disks with share interface can support high rates:  
33MB/s(Ultra-DMA), 40MB/s(SCSI-3), 400MB/s(FibreChannel)



CS 1555: Database Management Systems - Constantinos Costa

# Disk Failures

- Mean Time Between Failure (MTTF)
  - Vendor's claim: 30K hrs (3.6 yrs) – 1.2M hrs (136 yrs)
  - Practice: 1000 disks with 1.2M hrs MTTF then on average 1 disk will fail in 120 hrs.



CS 1555: Database Management Systems - Constantinos Costa

## But first, how is data stored on disk?



CS 1555: Database Management Systems - Constantinos Costa

## Data Elements

- **Field:** a database attribute (sequence of bytes)
- **Record:** sequence of fields

```
CREATE TABLE Student (
  Sid INTEGER,
  Name CHAR (24) ,
  Age INTEGER
);
```

0	3 4	27 28	31
Sid	Name	Age	



CS 1555: Database Management Systems - Constantinos Costa

## Data Elements

- **Field:** a database attribute (sequence of bytes)
- **Record:** sequence of fields

```
CREATE TABLE Student (
  Sid INTEGER,
  Name CHAR (24) ,
  Age INTEGER
);
```

0	3 4	27 28	31
Sid	Name	Age	

- **Block:** sequence of records
- **File:** sequence of blocks



CS 1555: Database Management Systems - Constantinos Costa

## Blocks & Files

```
CREATE TABLE Student (
  Sid INTEGER,
  Name CHAR (24) ,
  Age INTEGER
);
```

<i>SID</i>	<i>Name</i>	<i>Age</i>
546007	Peter	18
546100	Bob	19
546107	Ann	21
546207	Jane	20
546240	John	24
546350	Ben	18
546420	Suzy	27
546500	Peter	20



CS 1555: Database Management Systems - Constantinos Costa

## Blocks & Files

```
CREATE TABLE Student (
  Sid INTEGER,
  Name CHAR (24) ,
  Age INTEGER
);
```

	<i>SID</i>	<i>Name</i>	<i>Age</i>
Record 0	546007	Peter	18
Record 1	546100	Bob	19
Record 2	546107	Ann	21
Record 3	546207	Jane	20
Record 4	546240	John	24
Record 5	546350	Ben	18
Record 6	546420	Suzy	27
Record 7	546500	Peter	20

Block 0

Block 1

Block Header	Record 0	Record 1	Record 2	Record 3
--------------	----------	----------	----------	----------



CS 1555: Database Management Systems - Constantinos Costa



# Data on Disk

Address  
(offset)  
not  
stored

Stored  
Dats

000000)	05 06 7E 6E 6E 6E 08 79 - AE CE EE 08 88 7F 7F 7F
000010)	88 BD 7E 7E 7E 89 7E 6E - 6E 6E 08 9E 6E 6E 6E 79
000020)	04 79 AE CE EE 08 88 7F - 7F 7F 88 89 7E 7E 7E 89
000030)	FC FB 0F FB FC 08 9E 6E - 6E 6E 79 F9 EE EE EE 08
000040)	F9 EE EE EE 08 09 EE EE - EE 09 FF FF FF FF FF 89
000050)	6E 6E 6E 08 7E 6E 6E 6E - 08 88 7F 8F 7F 88 04 08
000060)	EF EF EF 08 FE FE 08 FE - FE FF 7E 08 7E FF 88 7F
000070)	8F 7F 88 06 BD 7E 7E 7E - 89 FF 7E 08 7E FF 7C 7A
000080)	6E 5E 3E 08 DF EF FB 08 - 7E 6E 6E 6E 08 BD 7E 7E
000090)	7E 89 00 00 00 00 00 00 - 00 00 00 00 00 00 00 00
0000A0)	00 00 00 00 00 00 00 00 - 00 00 00 00 00 00 00 00
0000B0)	00 00 00 00 00 00 00 00 - 00 00 00 00 00 00 00 00
0000C0)	00 00 00 00 00 00 00 00 - 00 00 00 00 00 00 00 00
0000D0)	00 00 00 00 00 00 00 00 - 00 00 00 00 00 00 00 00
0000E0)	00 00 00 00 00 00 00 00 - 00 00 00 00 00 00 00 00
0000F0)	00 00 00 00 00 00 00 00 - 00 00 00 00 00 00 00 00



CS 1555: Database Management Systems - Constantinos Costa

# Records on Disk

Address  
(offset)  
not  
stored

Stored  
Data

Block 0	05 06 7E 6E 6E 6E 08 79 - AE CE EE 08 88 7F 7F 7F	Record 0
	88 BD 7E 7E 7E 89 7E 6E - 6E 6E 08 9E 6E 6E 6E 79	
000020)	04 79 AE CE EE 08 88 7F - 7F 7F 88 89 7E 7E 7E 89	Record 1
000030)	FC FB 0F FB FC 08 9E 6E - 6E 6E 79 F9 EE EE EE 08	
000040)	F9 EE EE EE 08 09 EE EE - EE 09 FF FF FF FF FF 89	Record 2
000050)	6E 6E 6E 08 7E 6E 6E 6E - 08 88 7F 8F 7F 88 04 08	
000060)	EF EF EF 08 FE FE 08 FE - FE FF 7E 08 7E FF 88 7F	Record 3
000070)	8F 7F 88 06 BD 7E 7E 7E - 89 FF 7E 08 7E FF 7C 7A	
Block 1	6E 5E 3E 08 DF EF FB 08 - 7E 6E 6E 6E 08 BD 7E 7E	Record 4
	7E 89 00 00 00 00 00 00 - 00 00 00 00 00 00 00 00	
0000A0)	00 00 00 00 00 00 00 00 - 00 00 00 00 00 00 00 00	Record 5
0000B0)	00 00 00 00 00 00 00 00 - 00 00 00 00 00 00 00 00	
0000C0)	00 00 00 00 00 00 00 00 - 00 00 00 00 00 00 00 00	Record 6
0000D0)	00 00 00 00 00 00 00 00 - 00 00 00 00 00 00 00 00	
0000E0)	00 00 00 00 00 00 00 00 - 00 00 00 00 00 00 00 00	Record 7
0000F0)	00 00 00 00 00 00 00 00 - 00 00 00 00 00 00 00 00	



CS 1555: Database Management Systems - Constantinos Costa

## Improving Access Time ?

- What are the costs in I/O requests?
- I/O cost = *disk access* + *queuing delays*
- Disk access is reduced with *data organization*
- Queuing delays are reduced with *scheduling*



CS 1555: Database Management Systems - Constantinos Costa

## Improving Access Time

? Data access time could be improved by:

1. Scheduling
2. Block transfer
3. Buffering and Prefetching
4. Cylinder-based Organization
5. Multiple disks
6. Record Placement



CS 1555: Database Management Systems - Constantinos Costa

## Queuing Delay



CS 1555: Database Management Systems - Constantinos Costa

## Queuing Delay

- **Queuing Delay** occurs when:
  - multiple queries are executed at the same time, and
  - those queries access the same disk at the same time (i.e., send I/O requests)
- The disk head can only serve one request at a time!
- **Solution:**
  - Smart scheduling of serving disk requests
  - Parallel processing (multiple disks organizations)



CS 1555: Database Management Systems - Constantinos Costa

## Disk scheduling: Elevator Algorithm

- **Idea** : Schedule reading of requested blocks in the order in which they appear under the head
- **Elevator Algorithm (Scan)** - Works like an elevator
  - The arm moves towards **one direction** serving all requests on the visited tracks
  - Reverses direction when no more pending requests
- **Advantage**:
  - Reduces average access time for unpredictable requests
- **Problem**:
  - It is more effective in many disk-access requests
  - The benefit is not uniform among requests



CS 1555: Database Management Systems - Constantinos Costa

## Variations of Elevator Algorithm

- **Scan or Elevator Algorithm**
  - services both directions and reverses direction when reaches the edge of the disk
- **C-Scan or Circular Elevator Algorithm**
  - all requests are serviced in only one direction
  - more equal performance for all head positions
- **Look**
  - Scan but changes direction when no more pending requests
- **C-Look**
  - Combination of Look and C-Scan
- **FCSan**
  - prevents "arm stickiness" by using two queues (Q1 & Q2)
  - When Scan uses Q1 with old requests, new requests go to Q2



CS 1555: Database Management Systems - Constantinos Costa

## Improving Access Time

❓ Data access time could be improved by:

1. Scheduling
2. **Block transfer**
3. Buffering and Prefetching
4. Cylinder-based Organization
5. Multiple disks
6. Record Placement



CS 1555: Database Management Systems - Constantinos Costa

## Block Transfer

- To minimize access time: data is always transferred from/to disks by **blocks** of bytes
- A disk block is an adjacent sequence of sectors from a single track of one platter
- Block size typically ranges from 512 bytes to 4KBytes.



CS 1555: Database Management Systems - Constantinos Costa

## Improving Access Time

□ Data access time could be improved by:

1. Scheduling
2. Block transfer
3. Buffering and Prefetching
4. Cylinder-based Organization
5. Multiple disks
6. Record Placement



CS 1555: Database Management Systems - Constantinos Costa

## Buffering / Caching

- **Idea:** Keep as many blocks in main memory as possible to maximize the likelihood that a block of data needed by a transaction is in main memory

□ Might run out of memory space → replace blocks!

- Replacement Policies:

1. LRU (Least Recently Used)
2. MRU (Most Recently Used)
3. LFU (Least Frequently Used)
4. Dual-greedy (e.g., LRU & LFU)
5. ...



CS 1555: Database Management Systems - Constantinos Costa

## Prefetching or Double Buffering

- ❓ **Situation:** Needed data are known, but the timing of the request is data-dependent
- ❓ **Idea:** speed-up access by pre-loading needed data
- ❓ **Cons:**
  - requires extra main memory
  - no help if requests are random (unpredictable)



CS 1555: Database Management Systems - Constantinos Costa

## Improving Access Time


- o Data access time could be improved by:
  1. Scheduling
  2. Block transfer
  3. Buffering and Prefetching
  4. Cylinder-based Organization
  5. Multiple disks
  6. Record Placement



CS 1555: Database Management Systems - Constantinos Costa

## Cylinder-based Organization

- **Observation:** Data such as in relations is likely to be accessed together
- **Idea:** Store such data on the same cylinder
  - Blocks are **"Next"** to each other
  - Blocks on a track or on a cylinder effectively involve only the first seek time and first rotational latency
- **Advantage:**
  - Approaches the theoretical transfer rate
  - Excellent if access can be predicted in advance
  - Only one process/transaction is using the disk



CS 1555: Database Management Systems - Constantinos Costa

## Multiple Disks

- **Observation:**
  - Disk drives continue to become smaller and cheaper
- **Idea:**
  - Use multiple disks to support **parallel** access
    - Disks with independent heads connected to a single disk controller
  - 2X20 GB drives is faster than a single 40GB drive
    - But more expensive...
  - More efficient if they are kept busy!
  - It also enhances system reliability thru redundancy!!



CS 1555: Database Management Systems - Constantinos Costa



## Improving Access Time

? Data access time could be improved by:

1. Scheduling
2. Block transfer
3. Buffering and Prefetching
4. Cylinder-based Organization
5. Multiple disks
6. Record Placement



CS 1555: Database Management Systems - Constantinos Costa

## Multiple Disks: Organizations

- Data **partitioning** over several disks
  - **Pros:** increased access rate
  - **Problem:** if popular data is stored on same disk → request collisions (hot spots)
  - **Cons:** the cost of several small disks is greater than a single one with the same capacity
- **Mirror** disks: Disks hold identical copies
  - **Pros:** Doubles read rate (disk 1 **OR** 2)
  - Does not have the problem of request collisions
  - Does not slow down write requests (disk 1 **AND** 2)
  - **Cons:** Pay cost for two disks to get the storage of one



CS 1555: Database Management Systems - Constantinos Costa


# RAID Technology



CS 1555: Database Management Systems - Constantinos Costa

# RAID Technology

- A major advance in secondary storage technology is represented by the development of **RAID**
  - Redundant Arrays of Inexpensive (independent) Disks
- Acts as a single high-performance large disk
- **Data Striping:** distributes data transparently over multiple disks to make them appear as a single



CS 1555: Database Management Systems - Constantinos Costa


# Data Striping

- **Bit-level striping:** Split groups of bits over different disks
  - Example: split each bit of a bytes over **8 disks**
  - Bit number  $x$  is written on disk number  $x$
  - Every disk participates in every access
    - every access reads **8 times** as many data
- **Block-level** striping for blocks of a file
- **Sector-level** striping for sectors of a block
- **RAID Goals:**
  - Parallelize accesses so the **access time is reduced**
  - **Combining** Striping, Mirroring and Reliability → levels of RAID




CS 1555: Database Management Systems - Constantinos Costa

# Raid 0




CS 1555: Database Management Systems - Constantinos Costa

## Raid 0-1



CS 1555: Database Management Systems - Constantinos Costa

## Levels of RAID



CS 1555: Database Management Systems - Constantinos Costa

## Solid State Disk

- **Idea:** A SSD is essentially an array of NAND-flash memory devices which include a controller that emulates a hard disk drive
  - flash translation layer (FTL)
  - single-level cell (SLC) flash
- **Organization (flash):**
  - One or more dies (chips), each 8192 blocks
  - Dies are organized in planes (e.g., 4 planes of 2018 blocks)
  - Blocks is a set of pages 128—256KB (i.e., 32—64 pages)
  - Page is the unit of read and write (program), 512B—4KB
    - 128 bytes are used for metadata
  - Page or Block is the unit of storage system write



CS 1555: Database Management Systems - Constantinos Costa

## Constraints of Solid State Disks

- **Constraints:**
  - **Erase-Constraint** (or erase-before-write): A page cannot be over-written unless the block that contains the page has been erased
  - **Write-Constraint:** Pages must be written out sequentially within a block, from low to high addresses.
    - if  $p_i$  is written,  $p_{i-1}$  cannot be written until the block is erased
  - **Wear-Constraint:** The number of times a block can be erased is limited – typically 10,000 to 100,000 times.
    - Cleaning and wear-leveling schemes are important
  - **Asymmetry between operations:** Write is more expensive than read (~3 times).



CS 1555: Database Management Systems - Constantinos Costa

## Properties of Solid State Disks

- **Properties:**

- Random and Sequential Read Access with equal cost
- Faster than HDD
- parallel requests, interleaving, striping.. 100K – 300K I/Os
- Energy Efficient..  $\sim 25/750\mu\text{J}$  Read/Write,  $\sim 220\mu\text{J/s}$  Idle
- Resilient & reliable
  - Shock resistant and extreme temperature fluctuations
  - immune to strong magnetic fields
  - With good wear-leveling can last 10 years
- Small factor and silent (unless fan is used)



- **Problem:** Higher cost per megabyte of storage



- Capacity doubles every year

CS 1555: Database Management Systems - Constantinos Costa

## Improving Access Time

? Data access time could be improved by:

1. Scheduling
2. Block transfer
3. Buffering and Prefetching
4. Cylinder-based Organization
5. Multiple disks
6. Record Placement



CS 1555: Database Management Systems - Constantinos Costa

## Two Store Approaches



Row Stores



Column Stores



CS 1555: Database Management Systems - Constantinos Costa

## Column Stores

**SELECT SID, gender  
FROM Students;**


gender	Birthday	Name	SID	Address



CS 1555: Database Management Systems - Constantinos Costa

## Row Stores

- ❑ Store fields in one record contiguously on disk
- ❑ Use small (e.g., 4K) disk blocks



CS 1555: Database Management Systems - Constantinos Costa

## Record Blocking

- **Block header** includes:
  - Block ID
  - Role info such as data block, index root block, etc.
  - Timestamp
  - Links to other blocks of the same table/file
  - Free-list
- **Blocking factor  $bfr = \lfloor B/R \rfloor$**  : records per block
  - B = block size and R = Record size
- For spanned organization,  
 **$Bfr$**  is the average # of records per block




CS 1555: Database Management Systems - Constantinos Costa



## Fixed-Length Records

```
CREATE TABLE Student (
  SID      INTEGER
  name     CHAR(16),
  address  CHAR(130),
  gender   CHAR(1),
  birthdate DATE
)
```

- Fields are stored in sequence as the corresponding attributes are declared
- DATE**: 10-char string YYYY-MM-DD  
Fixed-length character string `char(10)`  
– example: 2002-09-15




CS 1555: Database Management Systems - Constantinos Costa

## Fixed-Length Records - Alignment

```
CREATE TABLE Student (
  SID      INTEGER,
  name     CHAR(16),
  address  CHAR(120),
  gender   CHAR(1),
  birthdate DATE
)
```

- Each record within a block starts at a byte that is multiple of 4
- Each field within a record starts at a byte off-set from the beginning of the record that is multiple of 4




CS 1555: Database Management Systems - Constantinos Costa

## Variable-Length Records

```
CREATE TABLE Student (
  SID      INTEGER,
  name     VARCHAR(16),
  address  CHAR(120),
  gender   CHAR(1),
  birthdate DATE
)
```

- Fields are stored in sequence as the corresponding attributes are declared
- DATE**: 10-char string YYYY-MM-DD  
Fixed-length character string `char(10)`  
– example: 2002-09-15




CS 1555: Database Management Systems - Constantinos Costa

## Variable-Length Attributes

- example: type **VARCHAR(16)**

– length + data:




– special end-of-filed symbol terminated:

- Any special character that does not appear in any field value,  
E.g., `¶`, `$`, `■`



– maximum length:



CS 1555: Database Management Systems - Constantinos Costa

## File Types

- Unordered files
- Ordered files
- Clustered files
- Hash files

*File header or descriptor includes:*

- field names and their data types
- address of the file block on disk



CS 1555: Database Management Systems - Constantinos Costa

## Unordered Files

- The simplest file structure: records are stored in no particular order
- Also called: **Heap**, Pile, or Random File
- New records are inserted at the **end** of file
  1. The last disk block is copied into buffer memory)
  2. New record is added
  3. Block is rewritten back to disk
- Record **insertion** is quite efficient



CS 1555: Database Management Systems - Constantinos Costa

## Example of Heap File

```
CREATE TABLE deposit (  
  account_number CHAR(10),  
  branch_name CHAR(22),  
  balance REAL  
)
```

Block 1	Record 0	A-102	Oakland	400
	Record 1	A-305	Shadyside	350
	Record 2	A-101	Downtown	700
	Record 3	A-222	Squirrel Hill	500
Block 2	Record 4	A-217	Shadyside	900
	Record 5	A-110	Waterfront	340
	Record 6	A-257	Oakland	600
	Record 7	A-403	Downtown	250



CS 1555: Database Management Systems - Constantinos Costa

## Properties of Unordered Files

- File records are inserted at the end of the file or in any file block with free space.
- Thus, insertion is efficient.
- To search for a record, a **linear search** through the file records is necessary which is quite expensive.
- Reading the records in order of any field requires sorting the file records.



CS 1555: Database Management Systems - Constantinos Costa

## Ordered Files

- ❑ Also called *sequential* files.
- ❑ File records are kept sorted by the value of an **ordering key** which has unique value (e.g., primary key)

```
CREATE TABLE deposit (  
  account_number CHAR(10),  
  branch_name    CHAR(22),  
  balance        REAL  
)
```

- Fixed-length records
- $10 + 22 + 8 = 40$  bytes

Record 0	A-101	Downtown	700
Record 1	A-102	Oakland	400
Record 2	A-205	Shadyside	350
Record 3	A-217	Shadyside	900
Record 4	A-222	Squirrel Hill	500
Record 5	A-310	Waterfront	340
Record 6	A-357	Oakland	600
Record 7	A-403	Downtown	250



CS 1555: Database Management Systems - Constantinos Costa

## Properties of Ordered Files

- Insertion is expensive: records must be inserted in the correct order.
- Deletion?
- Search for a record on its ordering field value is quite efficient (**binary search algorithm**).
- Search for a record on a non-ordering field?
- Reading the records in order of the ordering field is also quite efficient.
- Reading the records in any order?



CS 1555: Database Management Systems - Constantinos Costa

## Clustered Files

- Order files with ordering field which is **not a key**
- Ordering field has not a unique value

DEPTNUM	1	1	1	2	2	3	3	3	3	3	3	4	4
NAME													
SSN													
JOB													
BIRTHDATE													
SALARY													

5	5	5	5	6	6	6	6	6	8	8	8	8


- ☐ Properties?
- ☐ Purpose ?



CS 1555: Database Management Systems - Constantinos Costa

## Clustered File: Example 2

- ☐ For efficient insertion, deletion and search, each group of records with different ordering field is stored on separate block



CS 1555: Database Management Systems - Constantinos Costa

## Hash Files

- Also called **direct** files
- External hashing maps keys to disk blocks
  - Works similar to internal hashing
  - Static hashing
- Dynamic file expansion
  - Linear hashing
  - Extensible hashing



CS 1555: Database Management Systems - Constantinos Costa

## Static Hashing

- Hashing converts the key of a record into an address in which the record is stored.
- An external **hash** function maps the key to the relative address of a **bucket** in which the record is stored.
  - if a file is allocated **s** buckets, the hash function must convert a key **k** into the relative address of the block:
$$h(k) \in \{0, \dots, s-1\}$$
- A bucket is either one disk block or a cluster of contiguous blocks
- A table stored in the header of the file maps relative bucket numbers to disk block addresses.



CS 1555: Database Management Systems - Constantinos Costa

## Collision


- Insertion of a new record may lead to **collision**.
  - No space in  $B = h(k)$
- *Probing or conflict resolution*:
  - *Open Addressing (Rehashing)*: If bucket  $h(k)$  is full, use another hash function until a bucket with a free space is found.
    - E.g., *linear probing*
$$\alpha = h(\text{key}) = \text{key} \bmod s$$
$$rh(\alpha) = h(\alpha + 1)$$
    - Not a very good technique for databases (Why?)
  - *Chaining*: Use overflow buckets.



CS 1555: Database Management Systems - Constantinos Costa

## Handling of the Overflow Area

- Three approaches:
1. **Common** overflow area for all blocks in the file.
    - Each block has a pointer to its first record in the overflow area.
    - Records belonging to the same block are linked by pointers.



CS 1555: Database Management Systems - Constantinos Costa




## Handling of the Overflow Area

### 2. **Share** overflow area by some blocks

- e.g., residing on the same cylinder.

### 2. **Individual** overflow areas

- Each block has its own overflow
- Records are not linked



CS 1555: Database Management Systems - Constantinos Costa

## Hashing Functions

- A good hash functions must
  - 1) be computed efficiently
  - 2) minimize the number of collisions by spreading keys around the file as evenly and uniform as possible.
- Example of good functions
  - truncation
  - division:  $h(\text{key}) = \text{key} \bmod s$
  - Mid-square
  - Folding or partitioning
  - Other ad hoc methods
- Order preserving hash functions:
  - Maintain records in order of hash field values



CS 1555: Database Management Systems - Constantinos Costa

## Pros and Cons of Hashing

- Excellent performance for searching on **equality on the key** used for hashing (assuming low density).
- Records are not ordered (heap files).  
⇒ Any search other than on equality is very expensive (linear search or involves sorting).
- Prediction of total number of buckets is difficult.
  - allocate a large space.
  - estimate a "reasonable" size and periodically reorganize.



CS 1555: Database Management Systems - Constantinos Costa

## Dynamic Hashing Methods

- Allow the file size to change as records are added or deleted.
  - Linear Hashing
    - No additional structure
  - Extendible Hashing
  - Binary Hashing



CS 1555: Database Management Systems - Constantinos Costa

## Linear Hashing

- The file size grows linearly, bucket by bucket.
  - The file starts with  $s$  main buckets, where  $s$  is a power of 2, and  $k$  overflow buckets.
  - Buckets are numbered from 0 to  $s-1$   
=> initial hashing function  $h_0(key) = key \bmod s$ .
  - Collisions are handled thru **chaining**.
- We keep the following info:
  - $B_{last}$ : A pointer to the **current** last bucket.  
Initially,  $B_{last} = s-1$ .
  - $B_{split}$ : A pointer to the bucket that should be split next.  
Initially,  $B_{split} = 0$



CS 1555: Database Management Systems - Constantinos Costa

## Insertion of a Record (Method 1)

- ☐ if there is a collision, push tuple to an overflow bucket.
- ☐ if there are no overflow buckets available, proceed as follows until one becomes available:
  - A new bucket is appended at the end of the hash table and  
 $B_{last} = B_{last} + 1$
  - Records in  $B_{split}$  bucket are hashed again using  
 $h_1(key) = key \bmod (2s)$ ,  
-- these will either remain in  $B_{split}$  or stored in  $B_{last}$   
⇒ *this may free an overflow bucket.*
    - $B_{split}$  becomes  $B_{split} + 1$ .
- ☐ if  $B_{last} = 2s-1$   
set  $s = 2s$ ,  $B_{last} = s-1$ ,  $B_{split} = 0$ ,  $h_0(key) = h_1(key)$
- ☐ proceed as above.



CS 1555: Database Management Systems - Constantinos Costa

## Insertion of a Record (Method 2)

- ☐ if there is a collision, push tuple to an overflow bucket.
- ☐ if there are no overflow buckets available, proceed as follows until one becomes available:
  - A new bucket is appended at the end of the hash table (but  $B_{last}$  does not change as in Method 1)
  - Records in  $B_{split}$  bucket are hashed again using
$$h_1(key) = key \bmod (2s),$$
    - these will either remain in  $B_{split}$  or stored in  $B_{last}$
  - $\Rightarrow$  this may free an overflow bucket.
  - $B_{split}$  becomes  $B_{split} + 1$ .
- ☐ if  $B_{last} < B_{split}$   
set  $s = 2s$ ,  $B_{last} = s - 1$ ,  $B_{split} = 0$ ,  $h_0(key) = h_1(key)$
- ☐ proceed as above.



CS 1555: Database Management Systems - Constantinos Costa

## Search for a record

- Search for a Record
  - $I = h_0(key)$
  - if  $I < B_{split}$  then  $I = h_1(key)$
  - search the bucket whose hash value is  $I$  (and its overflow, if any).



CS 1555: Database Management Systems - Constantinos Costa

## Variations

- Split a bucket either when there is a need for an overflow bucket or based on some threshold  $u$ , e.g.,  $u > 0.9$
- $u$  is an upper bound of the *file load factor* or number of records that current buckets may store,

$$l = r / (bfr * N)$$

$r$  = current number of records,  $N$  = current number of buckets/

- Combination can be triggered when, e.g.,  $u < 0.7$
- Use a single hashing function:

$$A = k \bmod 2s$$

if  $A > B_{last}$

then  $A = A - s$

else  $A$



CS 1555: Database Management Systems - Constantinos Costa

## File Types

- Unordered files
- Ordered files
- Clustered files
- Hash files

**File header or descriptor includes:**

- field names and their data types
- address of the file block on disk



CS 1555: Database Management Systems - Constantinos Costa

## Dynamic Hashing Methods


- Allow the file size to change as records are added or deleted.
  - Linear Hashing
    - No additional structure
  - Extendible Hashing
  - Binary Hashing



CS 1555: Database Management Systems - Constantinos Costa

## Extendible Hashing

- The file is structured into two levels:  
***directory*** and buckets.



CS 1555: Database Management Systems - Constantinos Costa

## Extendible Hashing

- Directory contains only pointers to buckets (no keys).
- Use some hash function to generate a **pseudokey** of  $b$ -bits (typically  $b=32$ )
- Use the first (or last)  $d$  bits to find the offset of the bucket pointer in the directory.
- $d$  is called the (global) **directory depth**. It may be stored in the directory.
- In each bucket a local depth  $d'$  is stored indicating the most (or least) significant bits common to all keys in that bucket ( $d' \leq d$ ).



CS 1555: Database Management Systems - Constantinos Costa

## Extendible Hashing – File Growth

- No overflow buckets. New buckets are added one by one.
- When a bucket  $B$  with local depth  $d'$  overflows,  $B$  is split into two buckets and all keys are rehash using  $d'+1$  bits.
- If after a split,  $d' > d$ , double the size of the directory ( $d=d+1$ ) to accommodate the growth.



CS 1555: Database Management Systems - Constantinos Costa