

CS1555 Recitation 8 - Solution

Objective:

1. To get started with JDBC and demonstrate transaction concurrency control on PostgreSQL.
 2. To practice Views
-

PART 1:

Transactions are characterized by the ACID properties (Ref: Wikipedia):

- o Atomicity Each transaction is treated as a single “unit”.
- o Consistency Any data written to the database must be valid according to all defined rules.
- o Isolation Concurrent execution of transactions leaves the database in the same state that would have been obtained if the transactions were executed sequentially.
- o Durability Once a transaction has been committed, it will remain committed even in the case of a system failure.

Isolation is ensured by Concurrency Control, which synchronizes the execution of transactions to ensure each executes in an isolation manner.

By SQL standard and PostgreSQL, transactions can execute under different isolation levels:

Isolation Level	Dirty Read	Nonrepeatable Read	Phantom Read	Serialization Anomaly
Read uncommitted	Allowed, but not in PostgreSQL	Possible	Possible	Possible
Read committed	Not possible	Possible	Possible	Possible
Repeatable read	Not possible	Not possible	Allowed, but not in PostgreSQL	Possible
Serializable	Not possible	Not possible	Not possible	Not possible

PostgreSQL supports multi-version concurrency control. When an MVCC database needs to update an item of data, it will not overwrite the original data item with new data, but instead creates a newer version of the data item. Thus, there are multiple versions stored. Therefore, dirty reads will never happen, i.e. PostgreSQL's Read Uncommitted mode behaves like Read Committed.

Read Committed is the default isolation level in PostgreSQL. In read committed isolation level, two successive SELECT commands can see different data, even though they are within a single transaction.

The Repeatable Read isolation level only sees data committed before the transaction began. This is a stronger guarantee than is required by the SQL standard for this isolation level and prevents both nonrepeatable read and phantom read.

The Serializable isolation level provides the strictest transaction isolation. It works exactly the same as Repeatable Read except that it also prevents serialization anomalies.

(More details at: <https://www.postgresql.org/docs/current/transaction-iso.html>)

In addition to setting isolation levels, you may also place locks explicitly. There are two kinds of locks: (1) read lock or shared lock and (2) write lock or exclusive lock.

- **SHARED**

Shared locks allow multiple transactions to read data, but do not allow any transaction to change that data. Multiple transactions can hold shared locks simultaneously.

- **EXCLUSIVE**

An exclusive lock allows only one transaction to update a particular piece of data (insert, update, and delete). When one transaction has an exclusive lock on a row or table, no other lock of any type may be placed on it.

PART 2:

Before we start:

- Download `rec8db.sql`, `TranDemo1.java`, `TranDemo2.java` from course website.
- Download the PostgreSQL JDBC Driver from below link:
 - <https://jdbc.postgresql.org/download/postgresql-42.2.5.jar>
- Put above files into one folder
- Open DataGrip and 2 terminal windows.
 - In DataGrip, we run queries to keep track of changes in the database.
 - In the Terminal 1, we run `TranDemo1.java`
 - In the Terminal 2, we run `TranDemo2.java`

Example 0: Getting Started

- Edit the `TranDemo1.java` and `TranDemo2.java`, change the username and password to your username and password that you use to login to the PostgreSQL server.
- Compile the Java files using:
 - `javac -cp postgresql-42.2.5.jar TranDemo1.java`
 - `javac -cp postgresql-42.2.5.jar TranDemo2.java`
- Execute `rec8db.sql` in DataGrip.
- In Terminal 1, run the below command:
 - `java -cp postgresql-42.2.5.jar:. TranDemo1 0`
- Now read the demo source file to learn how it works. Note in the file:
 - How to connect to the DB.
 - How to execute an SQL statement.
 - How to iterate through the results set.

Notes:

- To run any of the examples, pass the example number as an argument.
- We will start running 2 transactions concurrently by running `TranDemo1` and `TranDemo2`.
 - Run `TranDemo1` first and then run `TranDemo2` while `TranDemo1` is still running.
 - Notice in the source codes how to group SQL statements into one transaction, commit/rollback the transaction and how to set isolation level for the transaction.
 - The sleep (milliseconds) function is used to force the statements in both transactions to execute in the order we want.

Example 1: Multi-version Concurrency of PostgreSQL

TranDemo1 (read committed)	TranDemo2 (read committed)
<pre>update class set max_num_students = 5 where classid = 1 sleep... rollback</pre>	<pre>SELECT * FROM class where classid = 1</pre>

Question: What is the max_num_students as read by TranDemo2?

Answer: Because PostgreSQL supports multi-version concurrency control, TranDemo2 read the committed value of max_num_students (i.e., before the update of TranDemo1). Therefore, TranDemo2 does not read dirty data and also does not have to wait for TranDemo1 to release the write lock (exclusive lock).

Example 2: (Implicit) Unrepeatable Read Problem

TranDemo1 (read committed)	TranDemo2 (read committed)
<pre>select max_num_students, cur_num_students from class where classid = 1 sleep... if (cur_num_students < max_num_students) update class set cur_num_students = cur_num_students + 1 where classid = 1 else print 'the class is full' commit</pre>	<pre>select max_num_students, cur_num_students from class where classid = 1 sleep... if (cur_num_students < max_num_students) update class set cur_num_students = cur_num_students + 1 where classid = 1 else print 'the class is full' commit</pre>

Question: What is the value of cur_num_students for class with classid = 1? Compare it to the max_num_classes.

Answer: Both of the two transactions registered for class 1, updating the cur_num_students to 3 even though the maximum number of students allowed in this class is only 2. The reason is, both of them read 1 as the current number of students in the class. This is called an implicit case of unrepeatable read: at the time TranDemo2 tried to update cur_num_students, it read that the value is still 1 while it has been updated to 2 (i.e., if it reads the value again at this point, the value will be different).

Example 3: Serializable Isolation Level:

The same as Example 2, but each transaction has the isolation level of **serializable**

Before running example 3, reset the database by rerunning rec8db.sql in the first terminal window.

Question: What is the value of cur_num_students for class with classid = 1 now? Do both transactions perform the update?

Answer: Only TranDemo1 can register (i.e. update cur_num_students,) successfully. TranDemo2 got the error message: “ERROR: could not serialize access due to concurrent update”. In serializable isolation level, PostgreSQL allows a transaction T to update a data item only if no other transaction is committing an update on that data item since T started. In this example, TranDemo1 has committed its update to the data item, preventing TranDemo2 from executing its update statement. Therefore, the data is still consistent: the cur_num_students is kept less than or equal to the max_num_students. The program should be able to catch this type of error from PostgreSQL and re-run the transaction instead of notifying the application user of the error.

Example 4: Using for Update of

The same as Example 2, but each transaction uses the following statement to select max and current number of students:

```
SELECT max_num_students, cur_num_students
FROM class where classid = 1
for update;
```

Again, before running example 4, reset the database by rerunning rec8db.sql in the first terminal window.

Question: What is the value of cur_num_students for class with classid = 1 now? Do both transactions perform the update?

Answer: Only TranDemo1 can register (i.e. update cur_num_students,) successfully. TranDemo2 got the user-friendly error message: “the class is full”. When a transaction executes a “select...for update” statement, it acquires an exclusive lock on the data item. This means that when TranDemo1 read the cur_num_students of class 1, it also kept an xlock on the corresponding row(s). Later, when TranDemo2 tried to read cur_num_students, it has to first ask for the xlock. Because TranDemo1 has the xlock, TranDemo2 has to wait until TranDemo1 commits and releases the lock. Therefore, the outcome of this example is the same as when the 2 transactions run sequentially.

Example 5: Deadlock

TranDemo1 (read committed)	TranDemo2 (read committed)
update class set max_num_students = 10 where classid = 1 sleep... update class set max_num_students = 10 where classid = 2 commit	update class set max_num_students = 20 where classid = 2 sleep... update class set max_num_students = 20 where classid = 1 commit

Question: What is the value of max_num_students? Do both transactions perform their updates?

Answer: Deadlock happens. One of the two transactions is selected as the victim and rollbacks. The victim transaction receives an error message that a deadlock is detected. The other transaction runs normally.

PART 3:

Before we start:

- Download and run the below to build the database:
 - Student_DB.sql

1. Create a view called student_courses that lists the SIDs, student names, number of courses in the Course_taken table.

```
create or replace view student_courses as
select s.sid, s.name, count(course_no) as num_courses
from student s, course_taken ct
where s.sid = ct.sid
group by s.sid, s.name;
```

2. Create a materialized view called mv_student_courses that lists the SIDs, student names, number of courses in the Course_taken table.

```
drop materialized if exists view mv_student_courses;
create materialized view mv_student_courses
as
select s.sid, s.name, count(course_no) as num_courses
from student s, course_taken ct
where s.sid = ct.sid
group by s.sid, s.name;
```

3. Execute the following commands. Compare the query results and time used of the two select statements.

```
insert into course_taken (course_no, sid, term, grade)
values ('CS1555', '129', 'Fall 18', null);
commit;

--REFRESH MATERIALIZED VIEW mv_student_courses;
select * from mv_student_courses;
select * from student_courses;
commit;
```

- The result from the materialized view is incorrect because the materialized view was not refreshed after the insert statement.
- The result from the view is correct because what a normal view does is rewriting the query. It does not store a snapshot of the query result like the materialized view.
- The running time of the materialized view is shorter, because it does not need to rewrite the query and run the rewritten query on the original Course_taken table.

4. Reset the database by running the Student_DB.sql and recreate the views. Comment back the line beginning with “REFRESH” in the above commands and execute the commands. Compare the query results of the two select statements.

- The user can request a refresh of the materialized view by running the command:
 - refresh materialized view <view_name>;
- The result from the materialized view is correct this time, because we refreshed the materialized view before the select statement.