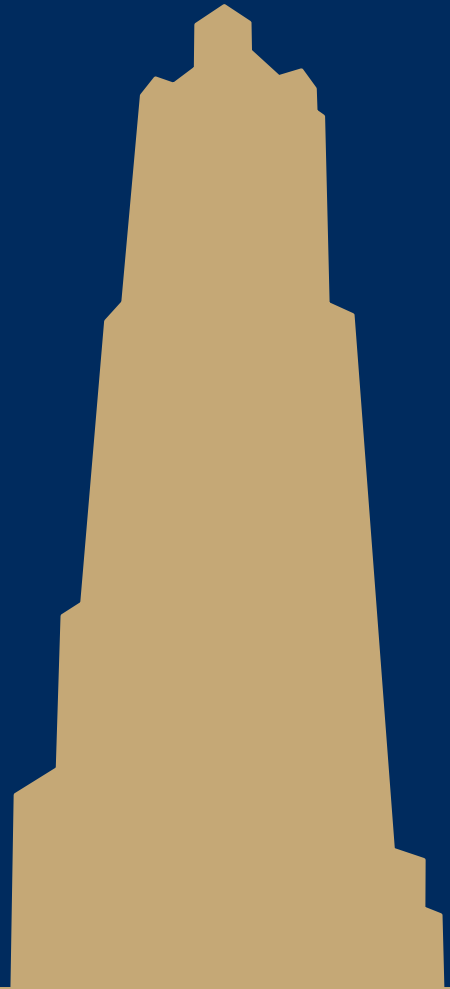


CS 1555

www.cs.pitt.edu/~nlf4/cs1555/

Using SQL as a DDL



SQL

- Declarative query language
- Originally developed for System R
 - First presented in 1974
- Current incarnation is the result of a very successful standardization effort by both ANSI and ISO
 - SQL-86 (SQL)
 - SQL-89 (SQL1)
 - SQL-92 (SQL2)
 - SQL:1999, SQL:2003, SQL:2006, SQL:2008, SQL:2011
- De-facto language for RDBMS
 - DDL, DML and VDL

The database schema

- Describes the data stored in the database
- Specifically:
 - Base relations
 - (tbl_name, creator, #tuples, tuple_length, #attributes...)
 - Attributes of relations (columns)
 - (tbl_name, attr_name, type, format, order, key_no, ...)
 - Indexes
 - (tbl_name, index_name, key_attribute,...)!
 - Authorizations
 - Integrity constraints
- All of this information is stored in the *catalog*
- SQL-92 and later allow multiple database schemas
 - Tables are named as SchemaName.TableName

Creating and deleting schemas

- **CREATE SCHEMA** *name* **AUTHORIZATION** *user*;
- **DROP SCHEMA** *name* [**RESTRICT** | **CASCADE**];
 - **RESTRICT**: removes schema if it doesn't contain any elements
 - **CASCADE**: remove schema and everything it contains

Creating database tables

- **CREATE TABLE** *tname* (*attribute list ...*);
 - Need to specify a name and a data type for each attribute

Numerical data types

- C-style integers
 - **SMALLINT**
 - **INT**
- Floating point types
 - **FLOAT**[(*p*)]
 - *p* sets the precision of the float, in bytes, max of 8
 - Allows custom precision floating point types
 - **REAL** and **DOUBLE PRECISION**
 - Set precision floating point types
 - 4 bytes and 8 bytes, respectively
- **DECIMAL**(*i*, *j*) or **NUMERIC**(*i*, *j*)
 - *i*: total number of digits (precision)
 - *j*: digits after a decimal point (scale)

Very few DBMSs "speak" standard SQL

- In reality, most will speak its own SQL dialect
 - Integer data type example:
 - PostgreSQL:
 - SMALLINT (16 bit)
 - INTEGER (32 bit)
 - BIGINT (64 bit)
 - My SQL:
 - SMALLINT (16 bit)
 - INT (32 bit)
 - BIGINT (64 bit)
 - Oracle:
 - SHORTINTEGER (16 bit)
 - INTEGER (32 bit)
 - LONGINTEGER (64 bit)
- Always consult the documentation for your DBMS!

Character string data types

- Strings of *printable* characters
- Enclosed in 'single quotes'
- **CHARACTER**(*n*) or **CHAR**(*n*)
 - Fixed length *n* strings
- **CHARACTER VARYING**(*n*) or **VARCHAR**(*n*)
 - Variable length string (max of *n*)
- Concatenation operator: **||**
 - 'abc' || 'XYZ' results in 'abcXYZ'
- **CLOB**(*size*)
 - Character Large Object
 - size specified in kilobytes (K), megabytes (M), or gigabytes (G)

Bit string data types

- Sequences of bits
- Enclosed in single quotes with a leading **B**
 - **B'001100101'**
- **BIT(*n*)**
 - Fixed length of *n* bit
- **BIT VARYING(*n*)**
 - Variable length (max of *n*)
- **BLOB(*size*)**
 - Binary Large Object
 - size specified in kilobytes (K), megabytes (M), or gigabytes (G)

Boolean values

- Valued **TRUE** or **FALSE**

NULL values

- Several reasons to store a **NULL** value in a tuple:
 - Unknown value
 - A person's date of birth is not known, so it is represented by **NULL** in the database.
 - Unavailable or withheld value
 - A person has a home phone but does not want it to be listed, so it is withheld and represented as **NULL** in the database.
 - Not applicable attribute
 - An attribute LastCollegeDegree would be **NULL** for a person who has no college degrees because it does not apply to that person.

Three-valued logic

- TRUE, FALSE, or UNKNOWN
- Consider:
 - Storing a NULL value for a BOOLEAN attribute
 - How should this be treated in a logical expression?
 - Evaluating a condition on a NULL value
 - Consider the partial condition:
 - Students.Name = 'SUSAN' AND Students.GPA > 2.0
 - If a row in the students table has a value of 3.0 for the GPA attribute and a NULL value for the Name attribute, should this condition be TRUE or FALSE?
 - What about NULL name, but a 1.0 for GPA?

Date and time data types

- **DATE**
 - E.g., YYYY-MM-DD
- **TIME [(p)] [WITH TIME ZONE]**
 - HH:MM:SS.dddddd{+,-}hhmm
 - If not specified, default is local timezone
- **TIMESTAMP [(p)] [WITH TIME ZONE]**
 - Complete date and time with up to 6 fractional seconds and optional time zone
- Dates and times must be valid!

Date and time implementations

- PostgreSQL sticks pretty close to SQL standard
- MySQL implements both TIMESTAMP and DATETIME
 - DATETIME is not a valid ANSI type
 - DATETIME range:
 - '1000-01-01 00:00:00' to '9999-12-31 23:59:59'
 - TIMESTAMP range in MySQL:
 - '1970-01-01 00:00:01' UTC to '2038-01-19 03:14:07' UTC
- Oracle DATE is not equivalent to ANSI DATE, it instead functions like ANSI TIMESTAMP

Intervals

- Represent periods of time and are used in operations on date and time data types
- `DATE {+,-} INTERVAL` results in `TIMESTAMP`
- `DATE - DATE` results in `INTERVAL`
- `INTERVAL {*,/} number` results in `INTERVAL`
- `INTERVAL {+,-} INTERVAL` results in `INTERVAL`
- Examples
 - `(CURRENT_DATE + INTERVAL '1 MONTH')`
 - `(CURRENT_DATE + INTERVAL '18 DAYS')`
 - `(CURRENT_DATE - BirthDate)`


Note

- Different DBMSs tend to use specialized DATE and TIME functions to manipulate DATE, TIME, and INTERVAL data
- Be sure to check the documentation of your DBMS

Interval specification

- **INTERVAL** [*fields*] [(*p*)]
- Possible fields:
 - YEAR
 - MONTH
 - DAY
 - HOUR
 - MINUTE
 - SECOND
 - YEAR TO MONTH
 - DAY TO HOUR
 - DAY TO MINUTE
 - DAY TO SECOND
 - HOUR TO MINUTE
 - HOUR TO SECOND
 - MINUTE TO SECOND
- Precision only applies to intervals with second field

All the data types you could want!

- DOMAIN defines datatype macros in a schema
 - Basic datatype
 - DEFAULT value
 - CHECK (validity conditions)
- **DROP DOMAIN *dname* [RESTRICT | CASCADE];**
 - RESTRICT drops domain if it is unused
 - CASCADE drops domain and replaces it with underlying type
- Examples:
 - **CREATE DOMAIN sectno_dom AS SMALLINT;**
 - **CREATE DOMAIN section_dom VARCHAR(20) DEFAULT 'none';**
 - **CREATE DOMAIN gpa_dom DECIMAL(3,2) DEFAULT 0.00;**
 - **CREATE DOMAIN valid_date DATE**
 **CONSTRAINT valid_date_range**
CHECK (VALUE BETWEEN '2017-01-01' AND '2018-01-01'));

Constraints on table attributes

- Constraints:
 - **NOT NULL**
 - **DEFAULT** *value*
 - without the **DEFAULT** clause, the default value is **NULL**
 - **PRIMARY KEY**(*attribute list*)
 - **UNIQUE**(*attribute list*)
 - allows the specification of alternative key
 - **FOREIGN KEY**(*key*) **REFERENCES** *table*(*key*)

CREATE TABLE example 1

```
CREATE TABLE Students
(
    ID            INTEGER,
    Name          VARCHAR(20),
    Major         VARCHAR(10),
    GPA           DECIMAL(3,2),
    CONSTRAINT Students_PK
        PRIMARY KEY (ID)
);
```

CREATE TABLE example 2

```
CREATE TABLE Enrollment
(
    Stud_ID INTEGER,
    Course VARCHAR(15),
    CONSTRAINT Enrollment_FK
        FOREIGN KEY (Stud_ID) REFERENCES Students(ID)
);
```

CREATE TABLE example 3

```
CREATE TABLE Students
(
    ID            INTEGER,
    Name          VARCHAR(20),
    Ssn           CHAR(9) NOT NULL,
    Major         VARCHAR(10),
    GPA           DECIMAL(3,2),
    CONSTRAINT Students_PK
        PRIMARY KEY (ID),
    CONSTRAINT Students_AK
        UNIQUE (Ssn)
);
```

Modifying tables after creation

- **ALTER TABLE *tname* ALTER COLUMN *cname* *options*;**
 - **ALTER TABLE Students ALTER COLUMN GPA DECIMAL(4,2);**
 - **ALTER TABLE Students ALTER COLUMN GPA SET DEFAULT NULL;**
 - **ALTER TABLE Students ALTER COLUMN GPA DROP DEFAULT;**
- **ALTER TABLE *tname* ADD COLUMN *cname* *type*;**
- **ALTER TABLE *tname* DROP COLUMN *cname* [RESTRICT|CASCADE];**
- **ALTER TABLE *tname* ADD CONSTRAINT *con_name* *description*;**
- **ALTER TABLE *tname* DROP CONSTRAINT *con_name*;**

Dropping tables

- **DROP TABLE *tname* [RESTRICT | CASCADE];**
 - RESTRICT drops the table only if it is not referenced
 - E.g., by constraints or views
 - CASCADE drops the table and items that reference it