

# Lecture 02: Relational Database Model

## CS 1555: Database Management Systems

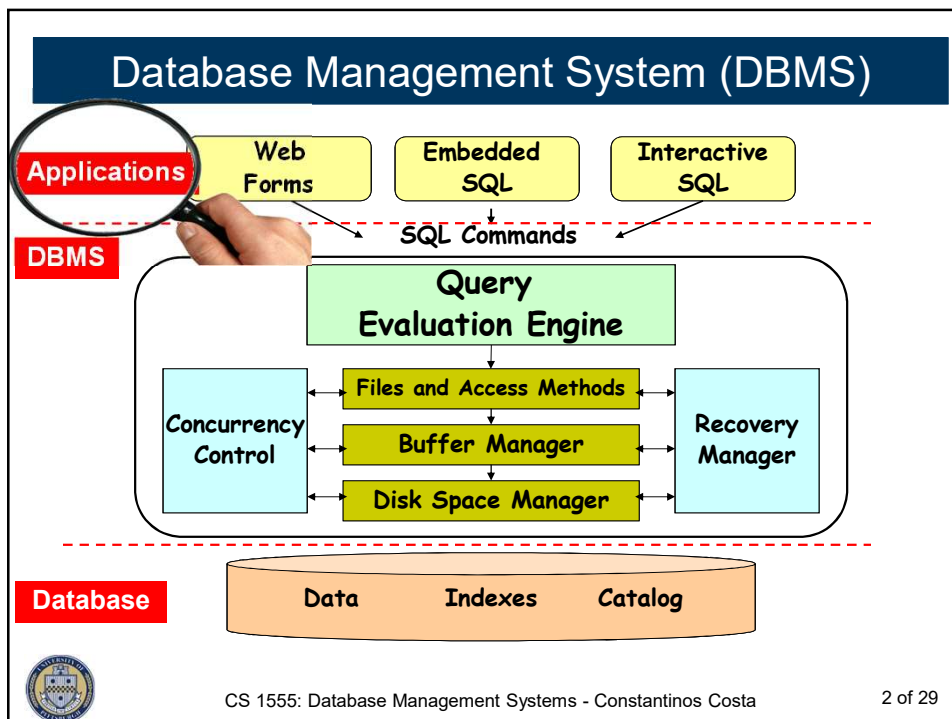
**Constantinos Costa**

[www.cs.pitt.edu/~nlf4/cs1555/](http://www.cs.pitt.edu/~nlf4/cs1555/) (tentative)

Jan 10, 2019, 16:00-17:15  
University of Pittsburgh, Pittsburgh, PA



Lectures based: P. Chrysanthis & N. Farnan Lectures



## The relational data model

- Proposed by E.F. “Ted” Codd in 1970
  - “A Relational Model of Data for Large Shared Data Banks.”
  - Built on the concept of the mathematical relation
  - Codd won a Turing award for this work in 1981
- First systems came about in 1977-1978
  - System-R and Ingres
- First commercial systems in the 1980’s
  - IBM and Oracle




CS 1555: Database Management Systems - Constantinos Costa

3 of 29

## Example of CODASYL Database

### Student Records



**FIND ANY <record-type> USING <fields list>**  
**GET {FIRST, NEXT, LAST} MEMBER WITHIN <set-type>**  
**WHERE <condition>**



CS 1555: Database Management Systems - Constantinos Costa

4 of 29

## Relational Model

<i>SID</i>	<i>Name</i>	<i>Major</i>	<i>GPA</i>
546007	Susan	CS	3.80
546100	Bob	CoE	3.65
546500	Bill	CS	3.70

**Students**

<i>CID</i>	<i>Name</i>
CS 1555	DB
CS 1530	SW
CS 1550	OS

**Courses**

<i>SID</i>	<i>CID</i>	<i>Grade</i>
546007	CS 1550	A
546007	CS 1530	B+
546100	CS 1550	B

**Enrollment**

- It is the most popular implementation model
  - Simplest, most uniform data structures, and is the most formal of all data model
- Both entity types and relationship types are represented by ***relations***, i.e., **tables**



CS 1555: Database Management Systems - Constantinos Costa

5 of 29

## The Mathematical Concept of Relation

- Let  $D_1, D_2, \dots, D_n$  be domains (not necessarily distinct), the Cartesian product of these  $n$  sets

$$D_1 \times D_2 \times \dots \times D_n$$

is the set of all possible ordered  $n$ -tuples

$$(v_1, v_2, \dots, v_n) \text{ such that } v_1 \in D_1, v_2 \in D_2, \dots, v_n \in D_n$$

- E.g., let  $D_1 = \{\text{Nick, Susan}\}$  and  $D_2 = \{\text{BS, MS, PhD}\}$   
 $D_1 \times D_2 = \{(\text{Nick, BS}), (\text{Nick, MS}), (\text{Nick, PhD}), (\text{Susan, BS}), (\text{Susan, MS}), (\text{Susan, PhD})\}$
- A relation is any subset of the Cartesian product
  - $R_1 = \{(\text{Nick, BS}), (\text{Nick, MS}), (\text{Susan, BS}), (\text{Susan, PhD})\}$
  - $R_2 = \{\}$



CS 1555: Database Management Systems - Constantinos Costa

6 of 29

## Two Notations

- Relation schema R is denoted by

$R = \{A_1:D_1, A_2:D_2, \dots, A_n:D_n\}$  or  $R = \{A_1, A_2, \dots, A_n\}$   
or  $R(A_1, A_2, \dots, A_n)$

– defines the name, the names of the attributes of that relation, and domains of those attributes

- Set-of-attributes

– A tuple t of r(R) is denoted by

$t = \{A_1:v_1, A_2:v_2, \dots, A_n:v_n\}, v_i \in D_i, 1 \leq i \leq n$  or

$t = \langle (A_1:v_1), (A_2:v_2), \dots, (A_n:v_n) \rangle, v_i \in D_i, 1 \leq i \leq n$

- List of attributes

– A tuple t of r(R) is denoted by

$t = (v_1, v_2, \dots, v_n), v_i \in D_i, 1 \leq i \leq n$



CS 1555: Database Management Systems - Constantinos Costa

7 of 29

## SQL Insert

- **Implicit (list):**  
**STUDENT(SID, Name, Major, GPA)**  
INSERT INTO STUDENT  
VALUES (165, 'Susan Jones', 'CS', 0.00);

- **Explicit (set):**  
INSERT INTO STUDENT (**SID, Name**)  
VALUES (165, 'Susan Jones');  
  
INSERT INTO STUDENT (**Name, SID**)  
VALUES ('Susan Jones', 165);

- Values-clause may be a list of tuples in some systems



CS 1555: Database Management Systems - Constantinos Costa

8 of 29

## Properties of Relations

- A relation is finite
- There are no duplicate tuples in a relation
  - Recall a relation is a set of tuples
- Order of tuples in a relation is not important
  - Many logical orders can be specified on a relation
- A value may appear multiple times in a column
- Order of attribute values in a tuple is
  - important in a *list-of-attributes* definition
  - not important in a *set-of-attributes* definition



CS 1555: Database Management Systems - Constantinos Costa

9 of 29

## Relation Schema

**STUDENT**

<i>SID</i>	<i>LName</i>	<i>Name</i>	<i>Class</i>	<i>Major</i>
123	Smith	John	3	CS
395	Aiken	Mary	4	CS

← Schema

### ◆ What is the meaning?

- A relation schema R specifies
  - The name of the relation
  - the attribute names  $A_i$  of R
  - the domain  $D_i$  (data type + format) for each attribute  $A_i$
- data type is a set of **atomic data** values:
  - no attribute is a set-valued (1st Normal Form, 1-NF)
  - no attribute is composite
- format specifies the representation of a data value



CS 1555: Database Management Systems - Constantinos Costa

10 of 29

## Example Table Schema

Schema of STUDENT(SID, Name, Major, GPA)

```
CREATE TABLE STUDENT
(  SID  INTEGER,
   Name CHAR(20),
   Major CHAR(4),
   GPA  DEC(3,2)
);
```



## Creating a Schema



- Corresponding database is at an empty state!
- **Initial state** when the database is populated (loaded)
- Domain (type) of each field is specified and enforced by the DBMS whenever tuples are added or modified



## Example: Domain Constraints

<i>SID</i>	<i>Name</i>	<i>Login</i>	<i>Age</i>	<i>GPA</i>
546007	Jones	jones@cs	18	3.4
546100	Smith	smith@ee	18	3.2
546500	Smith	smith@math	19	3.8

- Example of **IC Violation**:

**UPDATE** *Students S*

**SET** *S.Age* = 'Eighteen' ❌ ❌ ❌

**WHERE** *S.Name* = Jones;



## Useful Terms

- Cardinality of a relation  $r(R)$ : # of tuples in  $r(R)$  (denoted by  $|r(R)|$ )
- Arity or degree of  $r(R)$ : # of attributes in  $R$  (denoted by  $|R|$ )

$$|r(R)| = 3$$

$$|R| = 4$$

<i>SID</i>	<i>Degree</i>	<i>Major</i>	<i>Year</i>
123	BS	Math	1992
064	BA	History	1991
445	PhD	CS	1999

- ◆  $|r(R)| \geq 0$  And  $|R| > 0$
- ◆ **Cardinality** is property of a relation
- ◆ **Arity** is property of relation schema or a relation



# Relational Database Schema

- A **database schema** is a set of relation schemas and a set of **integrity constraints**




## □ Integrity Constraints

- **Structural** Integrity Constraints
  - **key** constraints: uniqueness of keys
  - **entity integrity** constraint: no primary key value can be **NULL**
  - **referential integrity** constraint
- **Semantic** Integrity Constraints
  - E.g., ??



# Integrity Constraints (ICs)

- **IC**: condition that must be true for *any* instance of the database (e.g., domain constraints)
  - A **legal** instance of a relation is one that satisfies all specified ICs
  - ICs are specified when schema is **defined**
  - ICs are enforced when tables are **modified**





## Example Table Schema in SQL (2)

Schema of STUDENT(SID, SSN, Name, Login, GPA)

```
CREATE TABLE STUDENT
(  SID  INTEGER NOT NULL,
  SSN   CHAR(9),
  Name  CHAR(20),
  Login CHAR(15),
  GPA   DEC(3,2),

  CONSTRAINT STUDENT_PK PRIMARY KEY (SID),
  CONSTRAINT STUDENT_UN_SSN
    UNIQUE (SSN)
  CONSTRAINT STUDENT_UN_Login
    UNIQUE (Login)
);
```



CS 1555: Database Management Systems - Constantinos Costa

17 of 29

## Example Table Schema in SQL (3)

Schema of STUDENT(SID, SSN, Name, Login, GPA)

```
CREATE TABLE STUDENT
(  SID  INTEGER
  CONSTRAINT STUDENT_PK_NOT_NULL NOT NULL
  CONSTRAINT STUDENT_PK PRIMARY KEY,
  SSN   CHAR(9)
  CONSTRAINT STUDENT_UN_SSN UNIQUE,
  Name  CHAR(20),
  Login CHAR(15),
  GPA   DEC(3,2),

  CONSTRAINT STUDENT_UN_Login
    UNIQUE (Login)
);
```



CS 1555: Database Management Systems - Constantinos Costa

18 of 29

## Identifying the Key

- What is the key in relation  
GRADUATE=(SID, Degree, Major, Year) ?

<i>SID</i>	<i>Degree</i>	<i>Major</i>	<i>Year</i>
123	BS	CS	1992
123	MS	CS	1993
064	BA	History	1991
445	PhD	CS	1999
123	BS	Math	1992
123	MS	Math	1992



## Foreign Keys

- Foreign key (FK) in relation  $R_2$  is a set of attributes of  $R_2$  that forms a primary key (PK) of another relation  $R_1$ .
  - Attributes in FK and PK have the same domain

### STUDENT

<u><i>SID</i></u>	<i>Degree</i>	<i>Major</i>	<i>Year</i>
-------------------	---------------	--------------	-------------

PK

### COURSE

<u><i>CID</i></u>	<i>Name</i>
-------------------	-------------

PK

### Enrolled

<u><i>SID</i></u>	<u><i>CID</i></u>	<i>Grade</i>
-------------------	-------------------	--------------

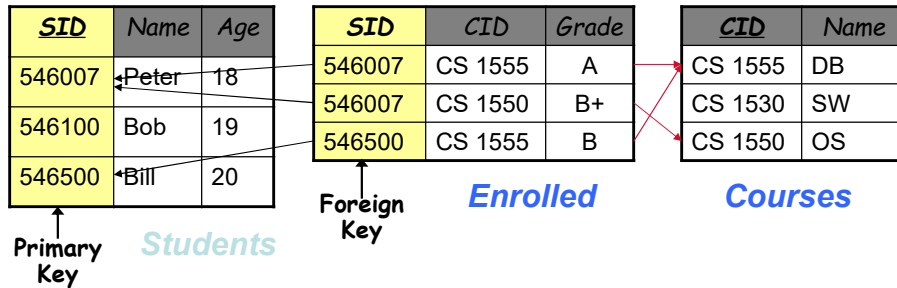
PK

FK

FK



## Foreign Key & Primary Key



- **Foreign key:** Set of fields in one relation that is used to “refer” to a tuple in another relation
  - Must correspond to primary key of the referred relation
  - E.g. *SID* is a foreign key referring to *Students*



CS 1555: Database Management Systems - Constantinos Costa

21 of 29

## Foreign Key Constraints

- If foreign key constraints are enforced, **referential integrity** is achieved
  - E.g.: Only students can enroll in a class
    - Only students listed in the “Students” relation should be allowed to enroll for courses
- Like a “*logical pointer*”
  - There shouldn’t be dangling references
    - Either valid PK or NULL



CS 1555: Database Management Systems - Constantinos Costa

22 of 29

## Any Attribute can be a Foreign Key

Faculty			Courses		
<u>FID</u>	Name	Area	<u>CID</u>	Name	Instructor
007	Panos	DB	CS 1555	DB	007
100	Daniel	OS	CS 1530	SW	NULL
500	Adriana	AI	CS 1550	OS	100

↑ Primary Key      ↑ Primary Key      ↑ Foreign Key

- **Foreign key:** Set of fields in one relation that is used to “refer” to a tuple in another relation
  - Must correspond to primary key of the referred relation
  - If not part of a key, it could be NULL



## Foreign Keys in SQL

<u>SID</u>	Name	Major	GPA	<u>CID</u>	Name	<u>SID</u>	<u>CID</u>	Grade
Students				Courses		Enrolled		

Red arrows indicate foreign key relationships: from SID in Enrolled to SID in Students, and from CID in Enrolled to CID in Courses.

```

CREATE TABLE Enrolled (
  SID CHAR(20), CID CHAR(20), Grade CHAR(2),
  CONSTRAINT Enrolled_PK PRIMARY KEY (SID, CID),
  CONSTRAINT Enrolled_FK_sid
    FOREIGN KEY (SID) REFERENCES Students (SID),
  CONSTRAINT Enrolled_FK_cid
    FOREIGN KEY (CID) REFERENCES Courses
);
  
```



## Referential Integrity Constraints

<i>SID</i>	<i>Name</i>	<i>Age</i>	<i>SID</i>	<i>CID</i>	<i>Grade</i>	<i>CID</i>	<i>CName</i>
546007	Peter	18	546007	CS 1555	A	CS 1555	DB
546100	Bob	19	546007	CS1530	B+	CS1530	SW
546500	Bill	20	546500	CS 1555	B	CS 1555	OS

*Students*                      *Enrolled*                      *Courses*

- Any IC Violation?

**DELETE**

**FROM** *Enrolled E*

**WHERE** *E.SID* = 546500;

**No Violations!** ✓



## Referential Integrity Enforcement

- What are the alternatives when a “Students” tuple is **deleted**?
  - Delete all** Enrolled tuples that refer to it
  - Disallow** deletion of a Students tuple that is referred to
  - Set** *SID* in Enrolled tuples that refer to it to some “**default**” *SID* (e.g., 000000)
  - If *SID* was not part of the primary key, **Set** *SID* to a special value “**NULL**”, denoting “*unknown*” or “*inapplicable*”




## Referential Integrity in SQL

- SQL/92 and SQL/99 support all 4 options on delete and update:
  - **NO ACTION** (default)
    - delete/update is rejected
  - **CASCADE**
    - also delete all tuples that refer to deleted tuple
  - **SET NULL / SET DEFAULT**
    - sets foreign key value of referencing tuple



## RI Trigger Actions in SQL



```
CREATE TABLE Enrolled (
    SID CHAR(20), CID CHAR(20), Grade CHAR(2),
    CONSTRAINT Enrolled_PK PRIMARY KEY (SID, CID),
    CONSTRAINT Enrolled_FK_sid
        FOREIGN KEY (SID) REFERENCES Students (SID),
    CONSTRAINT Enrolled_FK_cid
        FOREIGN KEY (CID) REFERENCES Courses
        ON UPDATE CASCADE ON DELETE NO ACTION
);
```



## Enforcing Integrity Constraints



- What would be the outcome?
  - Insert (585811, 'Jie', 19, 3.95) into Students
  - Insert (585811, NULL, NULL) into Enrollment
  - Insert (546100, 'CS 1555', NULL) into Enrollment
  - Insert (546100, 'Mary', 18, 3.65) into Students
  - Delete ('CS 1530') from Courses

<i>SID</i>	<i>Name</i>	<i>Age</i>	<i>GPA</i>
546007	Susan	18	3.8
546100	Bob	19	3.65
546500	Bill	20	3.7

**Students**

<i>CID</i>	<i>Name</i>
CS 1555	DB
CS 1530	SW
CS 1550	OS

**Courses**

<i>SID</i>	<i>CID</i>	<i>Grade</i>
546007	CS 1550	A
546007	CS 1530	B+
546100	CS 1550	B

**Enrollment**

