


CS1555/CS2055 Recitation 13

Objective: Practice recovery, operations on static hashing and extendible hashing


Part 1. Recovery

For the following transaction executions, state what the system should do when it restarts after a crash:

a. without checkpoints:



b. with checkpoints



Part 2. Hash Files


Hashing

- Convert the key of a record into an address in which the record is stored
- Search condition *must be* an equality condition on a single field, called **hash field**
- **Hash function:** a function that is applied to the hash field value of a record and yields the *address* of the disk block in which the record is stored

External Hashing for disk Files

- Target address space is made of **buckets**, each of which holds multiple records.
- A **bucket** is either one disk block or a number of contiguous disk blocks
- The *hashing function* maps a key into a relative bucket number, rather than assigning an absolute block address to the bucket

- A table maintained in the file header converts the bucket number into the address of the first disk block of the bucket
- Why target buckets and not blocks?



1) Static hashing:

- Converts the key of a record into an address in which the record is stored
- A fixed number of buckets, M , is allocated to match the fixed (static) Hash function range
- Each bucket has the same number of blocks. The number of blocks can change
- The collision problem is solved by using overflow buckets

Consider the following record keys: (3, 2, 1, 8, 6, 4, 14, 5, 9). Create the static hash structure, with $M=4$ main buckets, that will contain the provided records, using the chaining technique. Use $h(k) = k \bmod M$ as a hashing function. Each bucket can hold 2 records.

2) Dynamic hashing:

a) Extendible hashing:

- An array of 2^d bucket addresses is maintained, where d is called **global depth** of the directory
- The integer value corresponding to the first(high-order) d bits of a hash value is used as an index to the array to determine a directory entry and the address in that entry determines the bucket in which the corresponding records are stored
- A local depth d' –stored with each bucket –specifies the number of bits on which the bucket contents are based
- The value of d can be increased or decreased by one at a time, thus doubling or halving the number of entries in the directory array.
- **Doubling** is required when the bucket whose local depth d' , equal with the global depth d , overflows

Create an extendible hash structure for these record keys: (2, 3, 4, 8, 1, 12, 9, 7). Use the 4 bit binary representation of the keys (2=0010, 3=0011, 4=0100, 8=1000, 1=0001, 12=1100, 9=1001, 7=0111). Use a bfr=2.