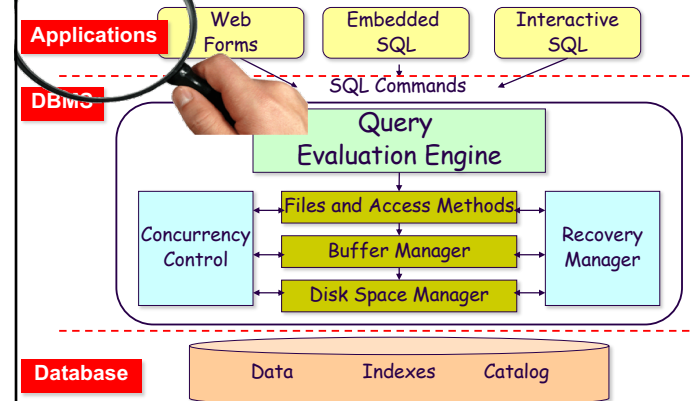


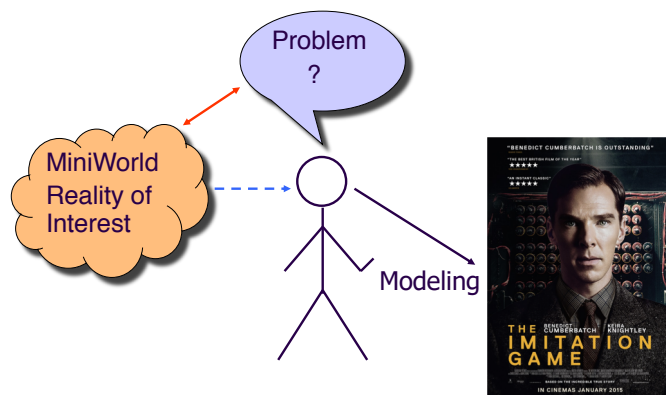
Relational Database Model

- ◆ Basic Concepts
- ◆ Mathematical Foundation

Database Management System (DBMS)



Data Modeling




Relational Model - History

- ❑ Before: records, pointers, sets, etc.
 - Hierarchical Data Model (IBM IMS, 1966-68)
 - Network Data Model (CODASYL DBTG, 1969)
- ❑ Introduced by E.F. Codd in 1970
- ❑ Revolutionary!
- ❑ First systems: 1977-8
 - System R; Ingres
- ❑ Turing award in 1981



Example of CODASYL Database

Student Records



FIND ANY <record-type> USING <fields list>
 GET {FIRST, NEXT, LAST} MEMBER WITHIN <set-type> WHERE <condition>

Relational Model

<i>SID</i>	<i>Name</i>	<i>Major</i>	<i>GPA</i>	<i>CID</i>	<i>Name</i>	<i>SID</i>	<i>CID</i>	<i>Grade</i>
546007	Susan	CS	3.80	CS 1555	DB	546007	CS 1555	A
546100	Bob	CoE	3.65	CS 1530	SW	546007	CS 1530	B+
546500	Bill	CS	3.70	CS 1550	OS	546100	CS 1550	B

Students

Courses

Enrollment

- It is the most popular implementation model
 - Simplest, most uniform data structures, and is the most formal of all data model
- Both entity types and relationship types are represented by **relations**, i.e., **tables**

The Mathematical Concept of Relation

- Let D_1, D_2, \dots, D_n be domains (not necessarily distinct), the Cartesian product of these n sets

$$D_1 \times D_2 \times \dots \times D_n$$
 is the set of all possible ordered n -tuples

$$(v_1, v_2, \dots, v_n)$$
 such that $v_1 \in D_1, v_2 \in D_2, \dots, v_n \in D_n$
- E.g., let $D_1 = \{\text{Nick, Susan}\}$ and $D_2 = \{\text{BS, MS, PhD}\}$

$$D_1 \times D_2 = \{(\text{Nick, BS}), (\text{Nick, MS}), (\text{Nick, PhD}), (\text{Susan, BS}), (\text{Susan, MS}), (\text{Susan, PhD})\}$$
- A relation is any subset of the Cartesian product
 - $R_1 = \{(\text{Nick, BS}), (\text{Nick, MS}), (\text{Susan, BS}), (\text{Susan, PhD})\}$
 - $R_2 = \{\}$

Two Notations

- Relation schema R is denoted by

$$R = \{A_1:D_1, A_2:D_2, \dots, A_n:D_n\} \text{ or } R = \{A_1, A_2, \dots, A_n\}$$
 or $R(A_1, A_2, \dots, A_n)$
- Set-of-attributes
 - A tuple t of $r(R)$ is denoted by

$$t = \{A_1:v_1, A_2:v_2, \dots, A_n:v_n\}, v_i \in D_i, 1 \leq i \leq n \text{ or }$$

$$t = \langle (A_1:v_1), (A_2:v_2), \dots, (A_n:v_n) \rangle, v_i \in D_i, 1 \leq i \leq n$$
- List of attributes
 - A tuple t of $r(R)$ is denoted by

$$t = (v_1, v_2, \dots, v_n), v_i \in D_i, 1 \leq i \leq n$$

SQL Insert

STUDENT(SID, Name, Major, GPA)

❑ Implicit (list):

```
INSERT INTO STUDENT
VALUES (165, 'Susan Jones', 'CS', 0.00);
```

❑ Explicit (set):

```
INSERT INTO STUDENT (SID, Name)
VALUES (165, 'Susan Jones');
```

```
INSERT INTO STUDENT (Name, SID)
VALUES ('Susan Jones', 165);
```

❑ Values-clause may be a list of tuples in some systems

Properties of Relations

- ❑ A relation is finite
- ❑ There are no duplicate tuples in a relation
 - Recall a relation is a set of tuples
- ❑ Order of tuples in a relation is not important
 - Many logical orders can be specified on a relation
- ❑ A value may appear multiple times in a column
- ❑ Order of attribute values in a tuple is
 - important in a *list-of-attributes* definition
 - not important in a *set-of-attributes* definition

Relation Schema

STUDENT	SID	LName	Name	Class	Major	Schema
	123	Smith	John	3	CS	
	395	Aiken	Mary	4	CS	

◆ What is the meaning?

- ❑ A **relation schema** R specifies
 - The name of the relation
 - the attribute names A_i of R
 - the domain D_i (data type + format) for each attribute A_i
- ❑ data type is a set of **atomic data** values:
 - no attribute is a set-valued (**1st Normal Form**, 1-NF)
 - no attribute is composite
- ❑ format specifies the representation of a data value

Example Table Schema

Schema of STUDENT(SID, Name, Major, GPA)

```
CREATE TABLE STUDENT
(
  SID    INTEGER,
  Name   CHAR(20),
  Major  CHAR(4),
  GPA    DEC(3,2)
);
```

Creating a Schema



- Corresponding database is at an **empty** state!
- Initial state** when the database is **populated** (loaded)
- Domain (type) of each field is **specified** and **enforced** by the DBMS whenever tuples are added or modified

Example: Domain Constraints

<i>SID</i>	<i>Name</i>	<i>Login</i>	<i>Age</i>	<i>GPA</i>
546007	Jones	jones@cs	18	3.4
546100	Smith	smith@ee	18	3.2
546500	Smith	smith@math	19	3.8

- Example of **IC Violation**:

```
UPDATE Students S
SET S.Age = 'Eighteen'   x x x
WHERE S.Name = Jones;
```

Useful Terms

- Cardinality** of a relation $r(R)$: # of tuples in $r(R)$ (denoted by $|r(R)|$)
- Arity** or **degree** of $r(R)$: # of attributes in R (denoted by $|R|$)

$$|R| = 4$$

$$|r(R)| = 3$$

<i>SID</i>	<i>Degree</i>	<i>Major</i>	<i>Year</i>
123	BS	Math	1992
064	BA	History	1991
445	PhD	CS	1999

- ◆ $|R| > 0$ and $|r(R)| \geq 0$
- ◆ Cardinality is property of a relation
- ◆ Arity is property of relation schema or a relation

Relational Database Schema

- A *database schema* is a set of relation schemas and a set of **integrity constraints**




- Integrity Constraints

- Structural** Integrity Constraints
 - **key** constraints: uniqueness of keys
 - **entity integrity** constraint: no primary key value can be **NULL**
 - **referential integrity** constraint
- Semantic** Integrity Constraints
 - E.g., ??

Integrity Constraints (ICs)

- ❑ **IC**: condition that must be true for *any* instance of the database (e.g., domain constraints)
 - A **legal** instance of a relation is one that satisfies all specified ICs
 - ICs are specified when schema is **defined**
 - ICs are enforced when tables are **modified**



Primary Key Constraint

- ❑ A set of fields is a **key** for a relation if :
 - No two distinct tuples can have same values in all key fields



- ❑ If there is more than one key for a relation:
 - Each is called a candidate key
 - One candidate key is designated as the **primary key**
 - Other candidate key(s) are designated as **alternative** or **unique key(s)**



Example of Keys

<i>SID</i>	<i>Name</i>	<i>Login</i>	<i>Age</i>	<i>GPA</i>
546007	Jones	jones@cs	18	3.4
546100	Smith	smith@ee	18	3.2
546500	Smith	smith@math	19	3.8

- ❑ **Candidate Keys**: *SID*, and *Login*
- ❑ **Primary Key**: *SID*
- ❑ **Unique Key**: *Login*

Example Table Schema in SQL

Schema of STUDENT(*SID*, *Login*, Name, Major, GPA)

```
CREATE TABLE STUDENT
(  SID  INTEGER NOT NULL,
   Login CHAR(15),
   Name  CHAR(20),
   Major CHAR(4),
   GPA   DEC(3,2),

   CONSTRAINT STUDENT_PK
      PRIMARY KEY (SID),
   CONSTRAINT STUDENT_UN
      UNIQUE (Login)      -- UNIQUE can take NULL values
);
```

Example Table Schema in SQL (2)

Schema of STUDENT(SID, Login, Name, SSN, GPA)

```
CREATE TABLE STUDENT
(
  SID INTEGER NOT NULL,
  Login CHAR(15),
  Name CHAR(20),
  SSN CHAR(9),
  GPA DEC(3,2),

  CONSTRAINT STUDENT_PK PRIMARY KEY (SID),
  CONSTRAINT STUDENT_UN_SSN
    UNIQUE (SSN),
  CONSTRAINT STUDENT_UN_Login
    UNIQUE (Login)
);
```

Example Table Schema in SQL (3)

Schema of STUDENT(SID, SSN, Name, Login, GPA)

```
CREATE TABLE STUDENT
(
  SID INTEGER
    CONSTRAINT STUDENT_PK_NOT_NULL NOT NULL
    CONSTRAINT STUDENT_PK PRIMARY KEY,
  SSN CHAR(9)
    CONSTRAINT STUDENT_UN_SSN UNIQUE,
  Name CHAR(20),
  Login CHAR(15),
  GPA DEC(3,2),

  CONSTRAINT STUDENT_UN_Login
    UNIQUE (Login)
);
```

Identifying the Key

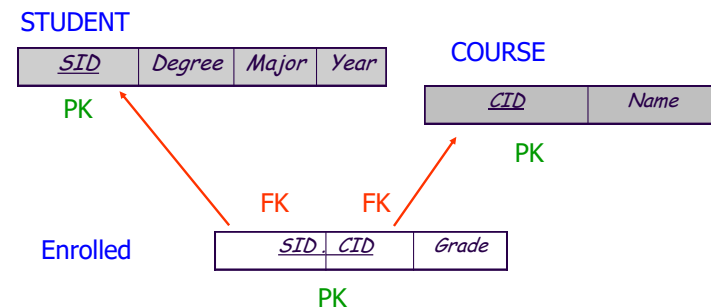
- What is the key in relation
GRADUATE=(SID, Degree, Major, Year) ?

<i>SID</i>	<i>Degree</i>	<i>Major</i>	<i>Year</i>
123	BS	CS	1992
123	MS	CS	1993
064	BA	History	1991
445	PhD	CS	1999
123	BS	Math	1992
123	MS	Math	1992




Foreign Keys

- Foreign key (FK)** in relation R_2 is a set of attributes of R_2 that forms a primary key (PK) of another relation R_1 .
 - Attributes in FK and PK have the **same domain**



Foreign Key & Primary Key




- Foreign key: Set of fields in one relation that is used to “refer” to a tuple in another relation
 - Must correspond to primary key of the referred relation
 - E.g. *SID* is a foreign key referring to *Students*

Foreign Key Constraints


- If foreign key constraints are enforced, **referential integrity** is achieved
 - E.g.: Only students can enroll in a class
 - Only students listed in the “Students” relation should be allowed to enroll for courses
- Like a “logical pointer”
 - There shouldn’t be dangling references
 - Either valid PK or NULL

Any Attribute can be a Foreign Key



- Foreign key: Set of fields in one relation that is used to “refer” to a tuple in another relation
 - Must correspond to primary key of the referred relation
 - If not part of a key, it could be NULL

Foreign Keys in SQL



```

CREATE TABLE Enrolled (
    SID CHAR(20), CID CHAR(20), Grade CHAR(2),
    CONSTRAINT Enrolled_PK PRIMARY KEY (SID, CID),
    CONSTRAINT Enrolled_FK_sid
        FOREIGN KEY (SID) REFERENCES Students (SID),
    CONSTRAINT Enrolled_FK_cid
        FOREIGN KEY (CID) REFERENCES Courses
);
    
```

Referential Integrity Constraints

<i>SID</i>	<i>Name</i>	<i>Age</i>	<i>SID</i>	<i>CID</i>	<i>Grade</i>	<i>CID</i>	<i>CName</i>
546007	Peter	18	546007	CS 1555	A	CS 1555	DB
546100	Bob	19	546007	CS1530	B+	CS1530	SW
546500	Bill	20	546500	CS 1555	B	CS 1550	OS

Students

Enrolled

Courses

- Any IC Violation?

DELETE

FROM *Enrolled E*

WHERE *E.SID* = 546500;

No Violations! ✓

Referential Integrity Constraints

<i>SID</i>	<i>Name</i>	<i>Age</i>	<i>SID</i>	<i>CID</i>	<i>Grade</i>	<i>CID</i>	<i>CName</i>
546007	Peter	18	546007	CS 1555	A	CS 1555	DB
546100	Bob	19	546007	CS1530	B+	CS1530	SW
546500	Bill	20	546500	CS 1555	B	CS 1550	OS

Students

Enrolled

Courses

- Any IC Violation?

DELETE

FROM *Students S*

WHERE *S.SID* = 546500;

Violation! ✗


Referential Integrity Enforcement

- What are the alternatives when a “Students” tuple is deleted?
 - Delete all** Enrolled tuples that refer to it
 - Disallow** deletion of a Students tuple that is referred to
 - Set** *SID* in Enrolled tuples that refer to it to some “default” *SID* (e.g., 000000)
 - If *SID* was not part of the primary key, **Set** *SID* to a special value “**NULL**”, denoting “unknown” or “inapplicable”

Referential Integrity in SQL

- SQL/92 and SQL/99 support all 4 options on delete and update:
 - NO ACTION** (default)
 - delete/update is rejected
 - CASCADE**
 - also delete all tuples that refer to deleted tuple
 - SET NULL / SET DEFAULT**
 - sets foreign key value of referencing tuple

RI Trigger Actions in SQL



```

CREATE TABLE Enrolled (
    SID CHAR(20), CID CHAR(20), Grade CHAR(2),
    CONSTRAINT Enrolled_PK PRIMARY KEY (SID, CID),
    CONSTRAINT Enrolled_FK_sid
        FOREIGN KEY (SID) REFERENCES Students (SID),
    CONSTRAINT Enrolled_FK_cid
        FOREIGN KEY (CID) REFERENCES Courses
        ON UPDATE CASCADE ON DELETE NO ACTION
);
    
```

Enforcing Integrity Constraints



- What would be the outcome?
 - Insert (585811, 'Jie', 19, 3.95) into Students
 - Insert (585811, NULL, NULL) into Enrollment
 - Insert (546100, 'CS 1555', NULL) into Enrollment
 - Insert (546100, 'Mary', 18, 3.65) into Students
 - Delete ('CS 1530') from Courses

<i>SID</i>	<i>Name</i>	<i>Age</i>	<i>GPA</i>
546007	Susan	18	3.8
546100	Bob	19	3.65
546500	Bill	20	3.7

Students

<i>CID</i>	<i>Name</i>
CS 1555	DB
CS 1530	SW
CS 1550	OS

Courses

<i>SID</i>	<i>CID</i>	<i>Grade</i>
546007	CS 1550	A
546007	CS 1530	B+
546100	CS 1550	B

Enrollment