

## Challenging yet Common Query

- Assume ENROLL(SID, CID, score)
- Find the ranking of students in CS2550 according to their scores. Your results should consider the case of tie.
  - E.g., output when 007 & 009 both received the highest score.

SID	Rank
007	1
009	1
003	3
005	4

## Nested Queries & Set Comparisons

- One of the most powerful features of SQL
- Two definitions of a set:
  - Explicit: list the members of the set within ( )  
E.g., (1, 2, 3) and ('Science', 'Art', 'Children')
  - Implicit: define it as a subquery (nested query) whose output table (set) can appear in any clause in place of a table
    - can be used at the FROM-clause
    - can be used at the WHERE-clause by the selection condition of the *outer* SELECT statement

## Scalar Subquery

- Scalar subquery: It is an inner query whose output is a single column and a single row
- This can be used in any expression in which a single value may appear
- E.g., Q1:

```
SELECT SID, Name
FROM STUDENT
WHERE SID = ( SELECT F.SID
               FROM STAFF AS F
               WHERE F.SSN = '132-32-2222');
```

## Another Example of Scalar Subquery

Q2: ?

```
SELECT S.SID, (SELECT DISTINCT MAX(Grade)
               FROM ENROLL E
               WHERE E.SID=S.SID) AS LG
FROM STUDENT S
WHERE S.Major = 'CS';
```

## Scalar Q2 Execution

### Students

SID	Name	Class	Major
123	John	3	CS
124	Mary	3	CS
126	Sam	2	CS
999	Newman	1	CS
132	Erin	2	EE

### Enroll

SID	CID	Term	Grade
123	CS1520	Fall 10	3.75
124	CS1520	Fall 10	4
126	CS1520	Fall 10	3
123	CS1555	Fall 10	4
124	CS1555	Fall 10	NULL
126	CS1550	Spring 11	NULL

### Q2 RESULT

SID	LG
123	4
124	4
126	3
999	

- Q2 is equivalent to a single-level SQL query; which one?

## Equivalent Query to Q2

Q2 Equivalent Query:

```
SELECT SID, MAX(Grade) AS LG
FROM STUDENT S NATURAL LEFT JOIN ENROLL E
WHERE S.Major = 'CS'
GROUP BY SID;
```

## Set Membership $x \in A$

- The IN and NOT IN operators check for simple membership.
- LIBRARIAN (SSN, Name, City, Gender, Salary, SNO);  
SECTION (SNO, Name, HeadSSN);

Q1: List each head librarian's SSN, along with their section, except those of the Science & Art sections.

```
SELECT HeadSSN, Name
FROM SECTION
WHERE Name NOT IN ('Science', 'Art');
```

## Set Membership

LIBRARIAN (SSN, Name, City, Gender, Salary, SNO);  
SECTION (SNO, Name, HeadSSN);  
DEPENDENT (Name, LIBSSN, DSSN, Gender, DoB);

- Q2: ?

```
SELECT L.Name
FROM LIBRARIAN AS L
WHERE L.SSN IN (SELECT LIBSSN
FROM DEPENDENT D
WHERE D.LIBSSN = L.SSN AND
D.Gender = L.Gender);
```

## Set Membership & Comparisons

- ❑ Test for membership on other comparisons:  
( =, <>, >, >=, <, <= )
- ❑ They can be quantified using **ANY** (i.e., SOME) or **ALL**.
- ❑ Q: List the SSN of all head librarians whose salary is lower than that of any librarian who is not a head librarian.

```
SELECT H.SSN
FROM LIBRARIAN H JOIN SECTION ON SSN = HeadSSN
WHERE H.Salary < ANY (SELECT L.Salary
                     FROM LIBRARIAN L
                     WHERE L.SSN NOT IN (SELECT HeadSSN
                                         FROM SECTION ));
```

## Set Comparisons: Unique & Empty

- ❑ The UNIQUE and NOT UNIQUE operators test for duplicates in a result (i.e, tests for a set and a bag, respectively)
- ❑ The EXISTS (not empty) and NOT EXISTS (empty) operators test for emptiness of a result

❑ Q: ?

```
SELECT L.SSN, L.Name
FROM LIBRARIAN L
WHERE NOT EXISTS (SELECT *
                  FROM SECTION
                  WHERE L.SSN = HeadSSN);
```

## Set Comparisons: Unique & Empty...

❑ Q: ?

```
SELECT S.SID
FROM STUDENT S
WHERE NOT UNIQUE ( SELECT *
                  FROM (SELECT SID
                        FROM STUDENT
                        WHERE Major = 'CS' ) UNION ALL
                  (SELECT SID
                   FROM STUDENT
                   WHERE Major = 'Math' )
                  WHERE S.SID = SID
                );
```

## Unique & Not Unique in Oracle

- ❑ UNIQUE in Oracle is used for constraints or instead of DISTINCT
  - Always use DISTINCT
- ❑ UNIQUE & Not UNIQUE can be expressed using JOIN

❑ Q: ?

```
SELECT C.SID
FROM Student C JOIN student M ON
    (C.SID = M.SID AND C.Major > M.Major);
```

## Removing Duplicates

- Q: List all the students who are double majors with their majors only once.
- Note that (1, CS, Math) and (1, Math, CS) are duplicates
- A: `SELECT C.SID, C.Major, M.Major`  
`FROM Student C JOIN student M ON`  
`(C.SID = M.SID AND C.Major > M.Major);`

Example:  
Student

SID	Major
123	CS
123	Math

Student C JOIN Student M

C.SID	C.Major	M.SID	M. Major
<del>123</del>	<del>CS</del>	<del>123</del>	<del>CS</del>
<del>123</del>	<del>CS</del>	<del>123</del>	<del>Math</del>
<b>123</b>	<b>Math</b>	123	<b>CS</b>
<del>123</del>	<del>Math</del>	<del>123</del>	<del>Math</del>

Result  
in bold

## Challenging yet Common Query

- Assume ENROLL(SID, CID, score)
- Find the ranking of students in CS2550 according to their scores. Your results should consider the case of tie.
  - E.g., output when 007 & 009 both received the highest score.

SID	Rank
007	1
009	1
003	3
005	4

## Ranking Query using of Scalar Subquery

- ENROLL(SID, CID, score)
- Q: Find the ranking of students in CS1555 according to their scores. Your results should consider the case of tie.

```
SELECT S.SID, (1 + (SELECT COUNT(*)
                    FROM ENROLL E
                    WHERE E.CID = 'CS 1555' AND
                          E.score > S.score)
              ) AS Rank
FROM ENROLL S
WHERE S.CID = 'CS 1555';
```

## Limiting result rows

- Save resources, speed-up result
- FETCH clause (SQL 2008)
  - OFFSET start { ROW | ROWS }
  - FETCH { FIRST | NEXT } [ count ] { ROW | ROWS } ONLY
  - the OFFSET clause must come before the FETCH clause
  - start: number of rows to skip [default 0]
  - count: maximum number of rows to return [default 1]
  - ROW, ROWS and NEXT is "noise" keywords – have no effect
- E.g., `SELECT *`  
`FROM STUDENT`  
`FETCH FIRST 10 ROWS ONLY;`

## Top-K Queries

❑ Q1: ?

```
SELECT *  
FROM STUDENT  
ORDER BY SID ASC  
FETCH FIRST 10 ROWS ONLY;
```

❑ Q2: ?

```
SELECT *  
FROM STUDENT  
ORDER BY QPA ASC  
FETCH FIRST 10 ROWS ONLY;
```

❑ How many tuples does Q2 return?

## Next-K Queries

❑ Q3: ?

```
SELECT *  
FROM STUDENT  
ORDER BY QPA ASC  
OFFSET 6 ROW  
FETCH NEXT 10 ROWS ONLY;
```

❑ Q3: (alternative that seems to work)

```
SELECT *  
FROM STUDENT  
ORDER BY QPA ASC  
OFFSET 6 ROW  
FETCH FIRST 10 ROWS ONLY;
```

## Limits (No Standard Syntax)

❑ `SELECT * FROM T WHERE ROWNUM <= 10;`

- Oracle (also supports the standard ?)

❑ `SELECT * FROM T LIMIT 10 OFFSET 20;`

- MySQL, PostgreSQL (supports the standard), SQLite

❑ `SELECT * FROM T`

`WHERE ID_T > 10 FETCH FIRST 10 ROWS ONLY;`

- IBM DB2

❑ `SELECT TOP 10 * FROM T;`

- MS SQL Server, Sybase ASE

## Top-K in Oracle

❑ Q: Select the top 3 students with the highest QPA

```
SELECT *  
FROM (SELECT * FROM STUDENT ORDER BY QPA DESC) S  
WHERE rownum <= 3  
ORDER BY rownum;
```

❑ ROWNUM: pseudocolumn which returns a number showing the order in which Oracle selects a row from a table

❑ Q: Select top students with offset = 3, next = 3

??

## Top-K in Oracle

- ❑ Q: Select top students with offset = 3, next = 3
- ❑ What about using **BETWEEN**?

```
SELECT *  
FROM (SELECT * FROM STUDENT ORDER BY QPA DESC)  
WHERE rownum BETWEEN 3 and 6;
```

Or

```
SELECT *  
FROM (SELECT * FROM STUDENT ORDER BY QPA DESC)  
WHERE rownum > 3 AND rownum <= 6;
```

- ❑ Does not work because **rownum** is assigned dynamically

## Next Top-K in Oracle

- ❑ Q: Select top students with offset = 3, next = 3


```
SELECT *  
FROM (  
    SELECT SID, Name, Major, QPA, rownum AS snum  
    FROM (SELECT * FROM STUDENT ORDER BY QPA DESC)  
    WHERE rownum <= 6)  
WHERE snum > 3;
```

- ❑ Consider QPAs: 1, 3, 6, 8, 2, 12, 19, 9

## Structured Query Language SQL – DML ◆ Update Operations

## Query Language with Update Statements?

- ❑ Is it a historical accident?
  - or Select is in the hard of updates!
  - *Selections* are expressed by the WHERE clause!



## Update Tuples

- Update can apply to a **single** relation
- Updates all the selected tuples by the condition in the WHERE-clause

- Examples:

```
UPDATE STUDENT
SET Name = 'Kathy Jones'
WHERE SID = 165;
```

```
UPDATE STUDENT
SET Major = 'CS'
WHERE DNO IN
( SELECT DNUM
  FROM DEPT
  WHERE Dname = 'CS' );
```

## Delete Tuple

- Delete removes all selected tuples by the condition in the WHERE-clause

- Examples:

```
DELETE FROM STUDENT
WHERE SID = 165;
```

```
DELETE FROM STUDENT
WHERE Name = 'John';
```

```
DELETE FROM STUDENT;
```

```
DELETE FROM STUDENT
WHERE DNO IN
( SELECT DNUM
  FROM DEPT
  WHERE Dname = 'CS' );
```

## Insert

**STUDENT(SID,Name,Major);**

- Two forms: **Implicit** (list) and **Explicit** (set)

- Implicit:**

```
INSERT INTO STUDENT
VALUES (165, 'Susan', 'CS');
```

- Explicit:**

```
INSERT INTO STUDENT (SID, Name)
VALUES (165, 'Susan Jones');
```

```
INSERT INTO STUDENT (Name, SID)
VALUES ( 'Susan Jones', 165);
```

## Derived Insert Values

- Tuples are derived using SELECT
- Useful to populate a table in the database from data already in the database

- E.g.,

```
INSERT INTO Dept_Info (Dept_Name, Num_Students)
SELECT Major, Count(*)
FROM STUDENT
GROUP BY Major;
```