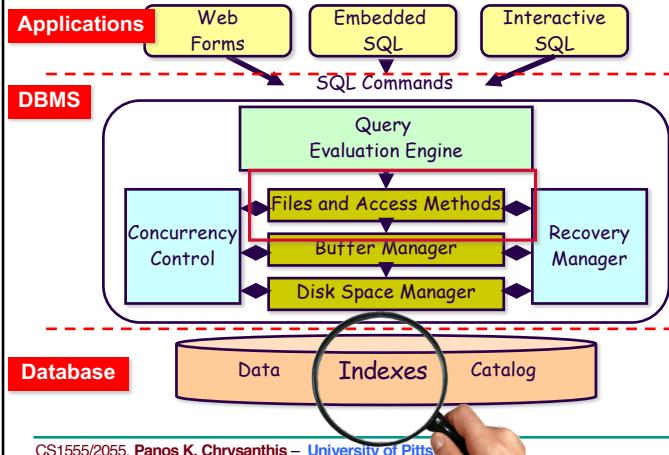


Access Paths or Index Structures for Files

CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

1

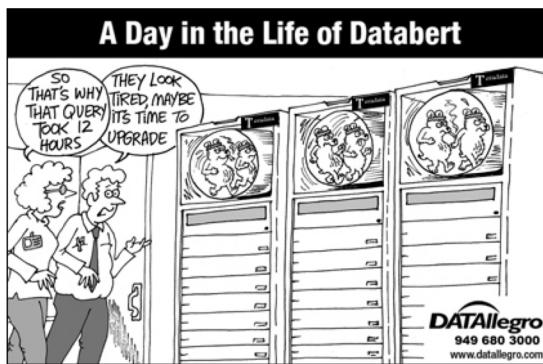
Database Management System (DBMS)



CS1555/2055, Panos K. Chrysanthis – University of Pitts

2

Queries are slow!



CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

3

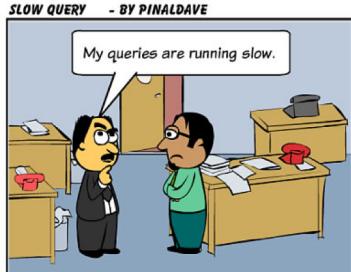
Create Index!



CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

4

Create Index!



CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

5

(C) SQLAuthority.com

Index Structures

- ❑ An index is an **auxiliary file** that makes it more efficient to search for a record in a **data file**.
- ❑ An index is usually specified on one field value of the data file.
- ❑ An index is an ordered file of entries
<field-value, pointer>
ordered by field value.
- ❑ Examples:
 - Index Sequential Access Method (ISAM)
 - primary, secondary and clustering indexes
 - B tree and B⁺ tree

CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

7

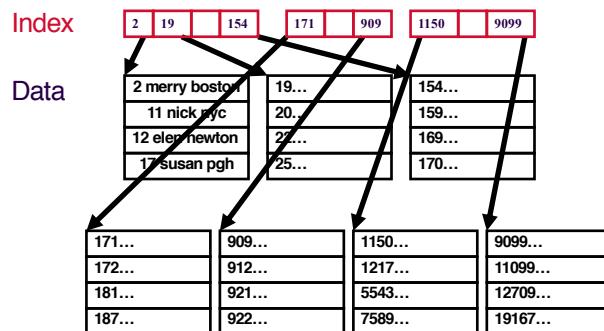
Index-Sequential Access Method (ISAM)

- ❑ It is a **primary index**.
- ❑ Defined on an **ordered** data file based on a key value.
- ❑ The first record in a data block is called **block anchor**.
- ❑ Includes one index entry **for each block** in the data file.
- ❑ The field value of an index entry is the key value of the block anchor.
- ❑ A similar scheme can use the **last record** in a data block.
- ❑ **Dense** index versus **Sparse** index.

CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

8

Example of Primary ISAM



CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

9

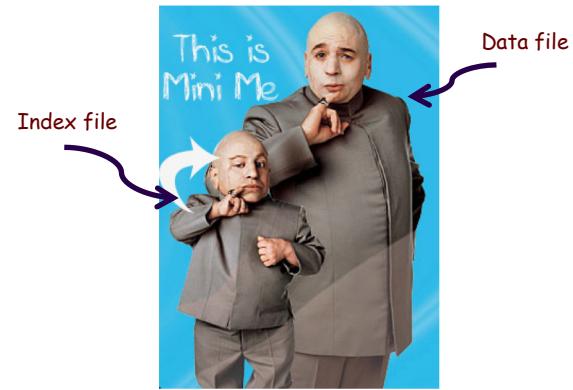
Advantages of a Primary Index

- ❑ A primary index might be stored in multiple blocks, but: it occupies much smaller space than a data file, because:
 1. There are fewer index entries than records
 2. Each index entry is typically smaller than a data record
 - Index record only 2 fields
- ❑ **More index entries than data records fit in a block**
- ❑ Binary search is more efficient on index file!
 - Let size of data file = B_{data} blocks
 - Let size of index file = B_{index} blocks
 - Typically: $B_{\text{index}} \ll B_{\text{data}}$ (much smaller)
 - Then: $\log_2 B_{\text{index}} < \log_2 B_{\text{data}}$

CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

10

Index File vs. Data File!



CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

11

Index Structure

- ❑ Single-level Indexes
 - Primary Indexes
 - Clustering Indexes
 - Secondary Indexes
- ❑ Multi-level Indexes

CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

14

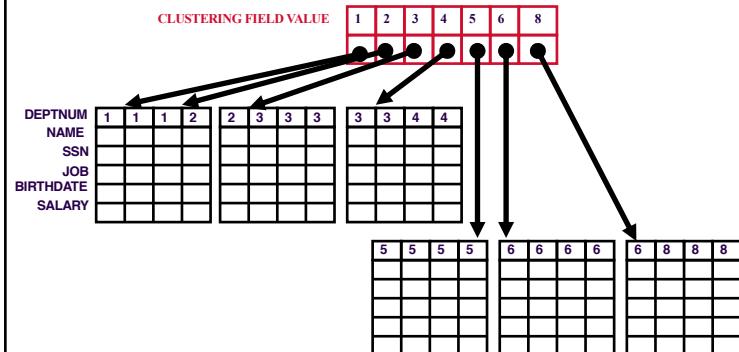
Clustering Index

- ❑ Defined on an **ordered** data file on a **non-key field**.
- ❑ One entry *for each distinct value* of the field
 - The index entry points to the **first data block** that contains records with that field value
- ❑ It is another example of **sparse** index

CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

15

Clustering Index: Example 1



CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

16

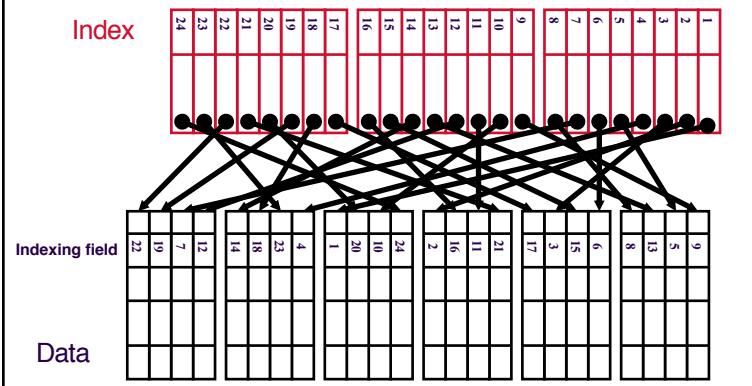
Secondary Index

- ❑ Also called ***nonclustering*** index.
- ❑ Defined on:
 - an **unordered** data file, or
 - on a **non-ordering** field of an **ordered** data file
- ❑ Can be defined on a key or non-key field.
- ❑ Includes one index entry for **each record** in the data file.
- ❑ Thus, it is a dense index and also called a **dense index**.
- ❑ The index entry points to ...?!

CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

18

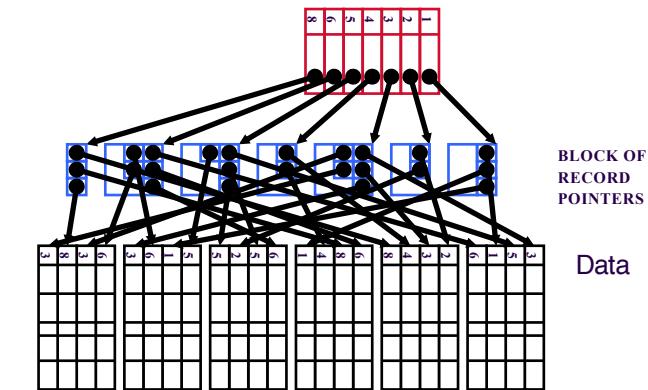
Secondary Index on key field



CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

19

Secondary Index on non-key field



CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

20

Summary

	Number of Entries	Dense/Sparse
Primary	?	?
Clustering	?	?
Secondary (key)	?	?
Secondary (nonkey)	?	?

Summary

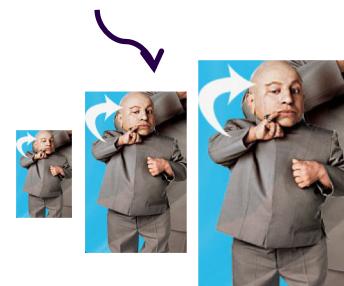
	Number of Entries	Dense/Sparse
Primary	Number of blocks in data file	Sparse
Clustering	Number of distinct values	Sparse
Secondary (key)	Number of records in data file	Dense
Secondary (nonkey)	Number of distinct values (1 st level)	Sparse
	Number of records/pointers (2 nd level)	Dense

Large Index

- ❑ For Big Databases
 - the Index could be very large to fit in main memory!
- ❑ *Can we do better?*

Multi-Level Index vs. Data!

Multi-level Index



Data file

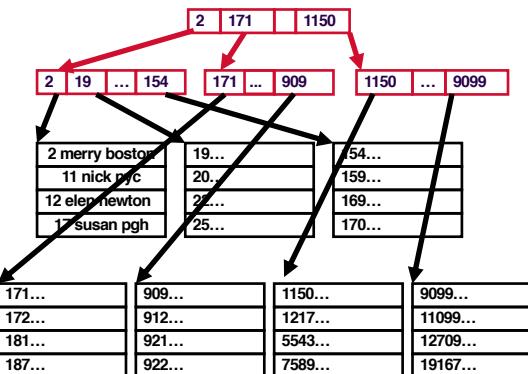
Multi-Level Indexes

- ❑ Because a single-level index is an **ordered file**: create a primary index to the index itself!
 - original index file is called **first-level index**
 - index to index is called the **second-level index**
- ❑ We can **repeat** the process until all entries of the top level fit in one disk block
- ❑ A multi-level index can be used on any type of first-level index: primary, secondary, clustering

CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

25

Multi-level ISAM



CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

26

Drawbacks of ISAM

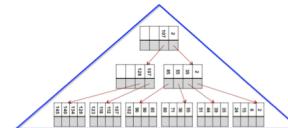
- ❑ A static structure.
 - It needs monitoring for dynamic databases.
- ❑ **Insertion/deletion** of new index entry is a problem, why?
 - every level of the index is an **ordered file**
 - Insertion is handled by some form of overflow blocks.
- ❑ Active files need frequent reorganization.
- ❑ No guaranteed performance for searching based on the key for active files.
 - anywhere between $O(\log n)$ to $O(n)$.

CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

29

Multi-level index as a tree

- ❑ Multi-level index is a form of search tree
 - Each node has pointers
 - By following a pointer, we restrict our search to a subtree and ignore all other nodes
 - The number of pointers (fan-out) equals the index blocking factor
- ❑ But the multi-level indexing we have seen so far is **static!**



CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

30

Is there:

One B+-tree to rule them all,
One B+-tree to find them,
One B+-tree to bring them all,
And in the darkness bind them

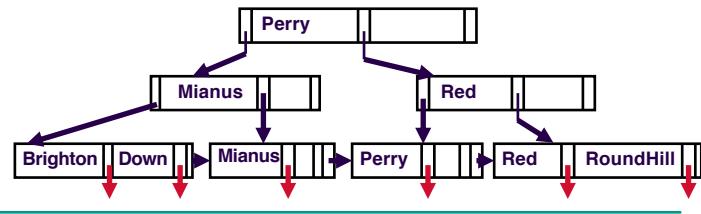


CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

31

B-trees and B⁺-trees

- ❑ Dynamic multi-level indexes.
- ❑ A multi-level index is a form of search tree
- ❑ B-trees and B⁺-trees are variations of search trees that allow efficient insertion and deletion of new values.
- ❑ Each node in the tree is a disk block
- ❑ Each node is kept between half-full and completely full.



CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

32

B/B⁺ tree Performance

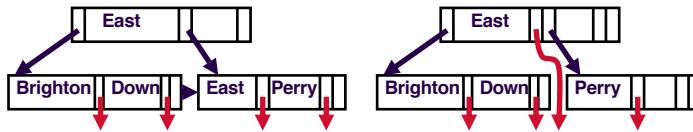
- ❑ Insertion:
 - Efficient if there is space
 - Otherwise a full node is split.
 - Splitting may propagate to other levels.
- ❑ Deletion:
 - Efficient if it does not cause the node to be less than half-full.
 - Otherwise it must be merged with its sibling node or, if not possible (i.e., sibling is full), accept half of the keys of its sibling node.

CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

33

B-trees Vs. B⁺-trees

- ❑ In a B⁺- tree, all pointers to data records exist at the leaf-level nodes.
- ❑ In a B-tree, pointers to data records exist at all levels.
- ❑ A B⁺- tree can have less levels than the corresponding B-tree.
- ❑ In a B⁺- tree, the search cost is the same for any value, $O(\log n)$. The tree is always balanced.



CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

34

B⁺-tree Index

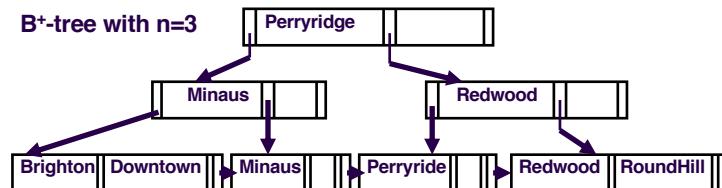
- ❑ A node is of the form:
 $[p_0, k_1, p_1, k_2, p_2, \dots, k_i, p_i, k_{i+1} \dots, k_n, p_n]$
- ❑ p_i 's are pointers and k_i 's are field values (keys)
- ❑ **Tree Order** is the number of pointers, e.g., n
- ❑ For every field value k in a node pointed to by p_i
 $k_i < k \leq k_{i+1}$ (alternative $k_i \leq k < k_{i+1}$)
- ❑ Every node, except for the root, has between $n/2$ and n children or pointers
 - internal: $\lceil n/2 \rceil$ tree pointers; leaf: $\lfloor n/2 \rfloor$ data pointers
- ❑ Leaf nodes are chain to form a link list (**fast sequential access**)

CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

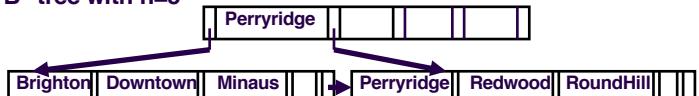
35

Examples of B⁺ trees

B⁺-tree with $n=3$



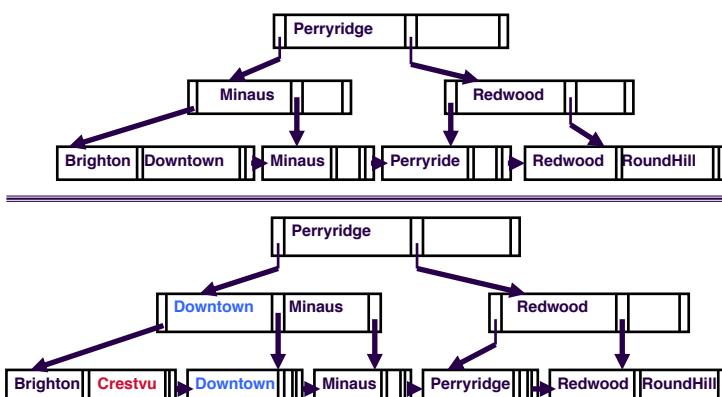
B⁺-tree with $n=5$



CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

36

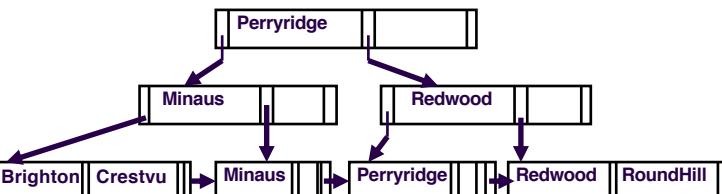
Insertion of “Crestvu”



CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

37

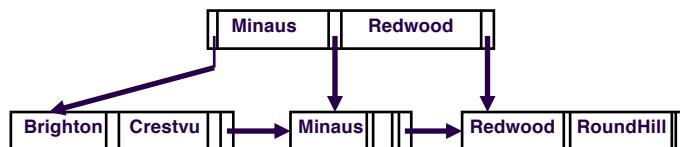
Deletion of “Downtown”



CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

38

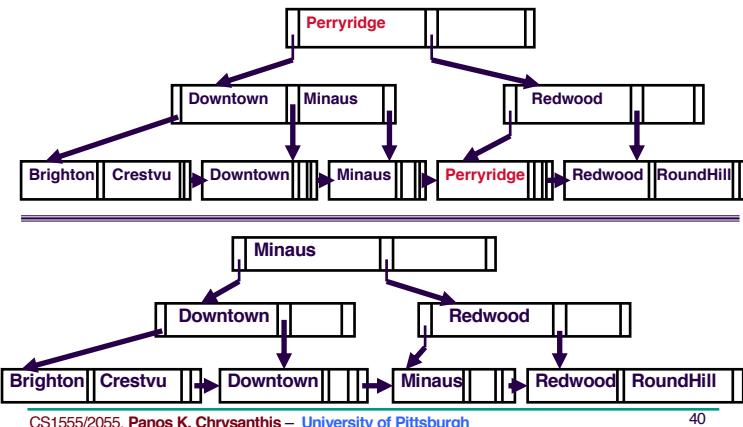
Deletion of “Perryridge”



CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

39

Deletion of “Perryridge” (w/out deletion of Downtown)



CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

40

Examples of B+ trees

B+-tree with n=4, i.e.,

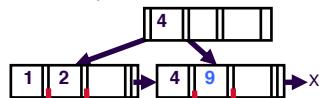
- internal nodes should have $\lceil 4/2 \rceil = 2$ tree pointers;
- Leaf nodes should have: $\lfloor 4/2 \rfloor = 2$ data pointers

Step 1: insert 1, 2, 4



Step 2: insert 9

Although you push 9 on the stack & inserted afterwards when allocating a new block, it is considered when deciding the split: 1 2 4. So the split is at 1 2 | 4 9.

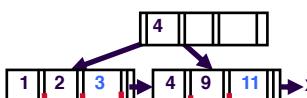


CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

41

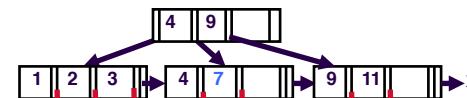
Examples of B+ trees

Step 3: insert 3, 11



Step 4: insert 7

Although you push 7 on the stack when allocating a new block, it is considered when deciding the split: 4 9 11. So the split is at 4 7 | 9 11

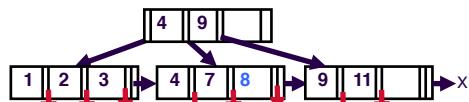


CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

42

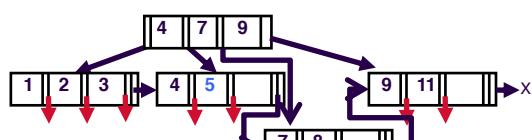
Examples of B+ trees

Step 5: insert 8



Step 6: insert 5

Although you push 5 on the stack when allocating a new block, it is considered when deciding the split: 4 7 8. So the split is at 4 5 | 7 8



Multiple-Key Access

Select name
From Student
Where
((State = 'PA') and
(year(Birthdate) = '1992'));

- Index only on State
- Index only on Birthdate
- Separate indexes on State and Birthdate:
 - (State or Birthday)
- Index on both State and Birthdate:
 - (State and Birthdate)

Indexes on Multiple Attributes

- Multiple Attribute** indexes use *composite search key*
 - Form of tuple of values: (a_1, a_2, \dots, a_n)
 - *Lexicographical ordering on tuples*
$$(a_1, a_2) < (b_1, b_2) \Rightarrow (a_1 < b_1) \vee ((a_1 = b_1) \wedge (a_2 < b_2))$$

- What about comparison conditions (range queries)?

Select name
From Student
Where ((State = 'PA') and
(year(Birthdate) < '1992'));

Or

Select name
From Student
Where ((year(Birthdate) < '1992') and
(State = 'PA'));

Point Access Methods (PAMs)

- Point Access Methods (PAMs)
 - A k-attribute record is envisioned as a point in a k-dimensional space
 - Can handle range queries
 - Can handle both points and spatial objects

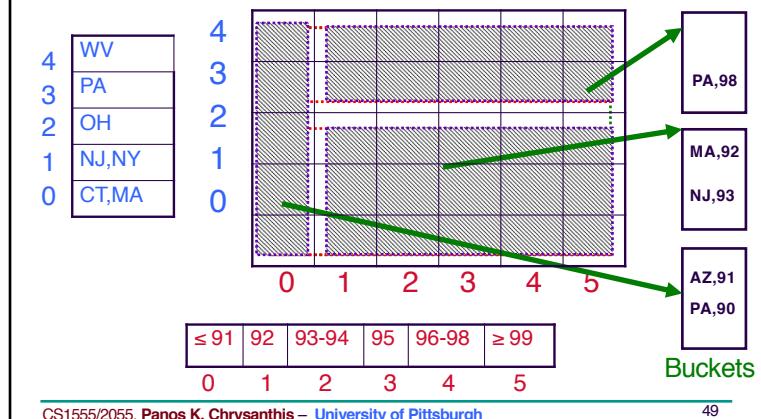
- Examples:

- Grid Files
- Quadtrees
- k-d trees
- R-trees

Grid Files

- ❑ Generalization of extensible hashing with a multi-dimensional directory, but static wrt directories
 - Fixed linear scales or dimensions/directories
 - Dynamic on bucket/block allocation
- ❑ Idea:
 - Impose a grid on the address space
 - Adapt grid to data density (re-organization)
 - Each grid cell corresponds to a disk block
 - One or more cells can share a disk block

Grid File (State, Year)



Grid Files Properties

- ❑ Pros:
 - Ensures 2 disk accesses for exact match
 - Symmetric w.r.t. the attributes
 - Adapts to non-uniform distributions
- ❑ Cons:
 - It does not work well if attributes are correlated
 - It requires extra space for directory which can grow large
 - If insertions are frequent, reorganization becomes costly
 - Hmm, how can this be addressed ?