

CS1555 Recitation 6 - Solution

Objective:

1. To practice top-k and ranking queries in PostgreSQL.
 2. To practice Views
-

PART 1:

Before we start:

- Download and run the below to build the database:
 - o Recitation_6_PostgreSQL.sql
 - o Recitation_6_Oracle.sql

1. List the student IDs and names of the 2 students with the lowest GPA.

-- Oracle and Postgres

```
select s.SID, s.name, avg(grade) as GPA
from course_taken ct
      join student s on ct.sid = s.sid
group by s.sid, s.name
order by GPA
FETCH FIRST 2 ROWS ONLY;
```

--or

```
select s.SID, s.name, avg(grade) as GPA
from course_taken ct
      join student s on ct.sid = s.sid
group by s.sid, s.name
order by GPA
limit 2;
```

--or

```
SELECT x.SID, x.name, x.GPA
FROM (
  select s.SID, s.name, avg(grade) as GPA,
  ROW_NUMBER() OVER (ORDER BY avg(grade))
  from course_taken ct
    join student s on ct.sid = s.sid
  group by s.sid, s.name) x
WHERE ROW_NUMBER <=2;
```

-- In oracle

```
select * from (
  select s.SID, s.name, avg(grade) as GPA
  from course_taken ct
    join student s on ct.sid = s.sid
  group by s.sid, s.name
  order by GPA)
WHERE rownum <= 2;
```

2. Rank the students (student ID and name) based on their GPA. Can we do something simpler?

```
select sid, name,
  (1 + (select count(*)
        from (select s.sid, s.name, avg(grade) as gpa
              from course_taken ct
                    join student s on ct.sid = s.sid
                    where grade is not null
                    group by s.sid, s.name
                    having avg(grade) > i.gpa
                    order by gpa) as e)
  ) as rank
from (select s.sid, s.name, avg(grade) as gpa
      from course_taken ct
            join student s on ct.sid = s.sid
            where grade is not null
            group by s.sid, s.name
            order by gpa) as i
order by rank;
```

-- Simplify

```
create or replace view student_gpa as
select s.sid, s.name, avg(grade) as gpa
from course_taken ct
      join student s on ct.sid = s.sid
where grade is not null
group by s.sid, s.name
order by gpa;
```

-- Now the query

```
select sid, name,
  (1 + (select count(*)
        from student_gpa as e
        where e.gpa > i.gpa)
  ) as rank
from student_gpa as i
order by rank;
```

PART 2:

Before we start:

- Download and run the below to build the database:
 - Recitation_6_PostgreSQL.sql
 - Recitation_6_Oracle.sql

1. Create a view called student_courses that lists the SIDs, student names, number of courses in the Course_taken table.

```
create or replace view student_courses as
select s.sid, s.name, count(course_no) as num_courses
from student s, course_taken ct
where s.sid = ct.sid
group by s.sid, s.name;
```

2. Create a materialized view called mv_student_courses that lists the SIDs, student names, number of courses in the Course_taken table.

```
drop materialized view if exists mv_student_courses;
create materialized view mv_student_courses
as
select s.sid, s.name, count(course_no) as num_courses
from student s, course_taken ct
where s.sid = ct.sid
group by s.sid, s.name;
```

3. Execute the following commands. Compare the query results and time used of the two select statements.

```
insert into course_taken (course_no, sid, term, grade)
values ('CS1555', '129', 'Fall 19', null);
commit;
```

```
--REFRESH MATERIALIZED VIEW mv_student_courses;
select * from mv_student_courses;
select * from student_courses;
commit;
```

- The result from the materialized view is incorrect because the materialized view was not refreshed after the insert statement.

- The result from the view is correct because what a normal view does is rewriting the query. It does not store a snapshot of the query result like the materialized view.
- The running time of the materialized view is shorter, because it does not need to rewrite the query and run the rewritten query on the original `Course_taken` table.

4. Reset the database by running the `Recitation_6_PostgreSQL.sql` and recreate the views. Comment back the line beginning with “REFRESH” in the above commands and execute the commands. Compare the query results of the two select statements.

- The user can request a refresh of the materialized view by running the command:
 - `refresh materialized view <view_name>;`
- The result from the materialized view is correct this time, because we refreshed the materialized view before the select statement.