

Relational Database Model

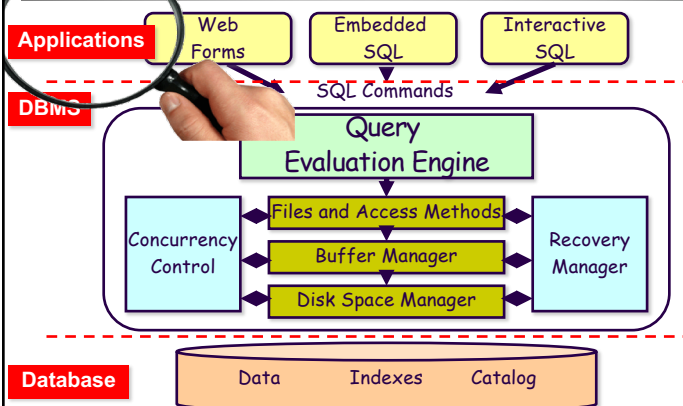
- ◆ Basic Concepts
- ◆ Mathematical Foundation

Performance Requirements

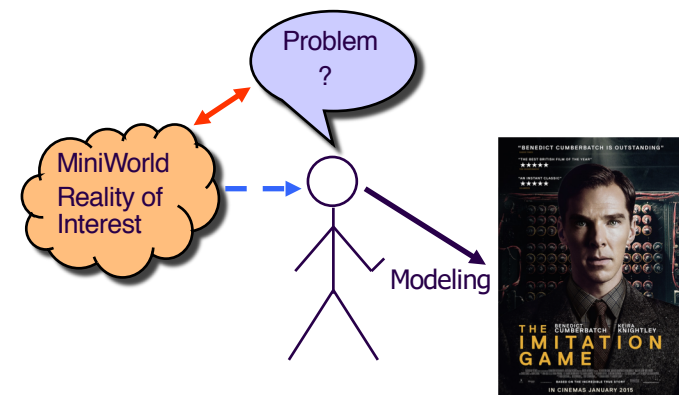
- Abstraction
 - Data abstraction
 - Execution abstraction
- Reliability
 - High availability: recovery time is *short*
 - Trusted/Quality data
- Efficiency/Performance
 - High throughput
(Committed transactions per unit time)
 - Short or bounded response time
 - Energy Efficiency



Database Management System (DBMS)



Data Modeling



Relationships




Relational Model - History

- ❑ Before: records, pointers, sets, etc.
 - Hierarchical Data Model (IBM IMS, 1966-68)
 - Network Data Model (CODASYL DBTG, 1969)
- ❑ Introduced by E.F. Codd in 1970
- ❑ Revolutionary!
- ❑ First systems: 1977-8
 - System R; Ingres
- ❑ Turing award in 1981



Example of CODASYL Database

Student Records



FIND ANY <record-type> USING <fields list>
 GET {FIRST, NEXT, LAST} MEMBER WITHIN <set-type> WHERE <condition>

Relational Model

<i>SID</i>	<i>Name</i>	<i>Major</i>	<i>GPA</i>	<i>CID</i>	<i>Name</i>	<i>SID</i>	<i>CID</i>	<i>Grade</i>
546007	Susan	CS	3.80	CS1555	DB	546007	CS1550	A
546100	Bob	CoE	3.65	CS1530	SW	546007	CS1530	B+
546500	Bill	CS	3.70	CS1550	OS	546100	CS1550	B

Students

Courses

Enrollment

- ❑ It is the most popular implementation model
 - Simplest, most uniform data structures, and is the most formal of all data model
- ❑ Both entity types and relationship types are represented by **relations**, i.e., **tables**

The Mathematical Concept of Relation

- Let D_1, D_2, \dots, D_n be domains (not necessarily distinct), the Cartesian product of these n sets
 $D_1 \times D_2 \times \dots \times D_n$
is the set of all possible ordered n -tuples
 (v_1, v_2, \dots, v_n) such that $v_1 \in D_1, v_2 \in D_2, \dots, v_n \in D_n$
- E.g., let $D_1 = \{\text{Nick, Susan}\}$ and $D_2 = \{\text{BS, MS, PhD}\}$
 $D_1 \times D_2 = \{(\text{Nick, BS}), (\text{Nick, MS}), (\text{Nick, PhD}), (\text{Susan, BS}), (\text{Susan, MS}), (\text{Susan, PhD})\}$
- A relation is any subset of the Cartesian product
 - $R_1 = \{(\text{Nick, BS}), (\text{Nick, MS}), (\text{Susan, BS}), (\text{Susan, PhD})\}$
 - $R_2 = \{\}$

Two Notations

- Relation schema R is denoted by
 $R = \{A_1:D_1, A_2:D_2, \dots, A_n:D_n\}$ or $R = \{A_1, A_2, \dots, A_n\}$
or $R(A_1, A_2, \dots, A_n)$
- Set-of-attributes
 - A tuple t of $r(R)$ is denoted by
 $t = \{A_1:v_1, A_2:v_2, \dots, A_n:v_n\}, v_i \in D_i, 1 \leq i \leq n$ or
 $t = \langle (A_1:v_1), (A_2:v_2), \dots, (A_n:v_n) \rangle, v_i \in D_i, 1 \leq i \leq n$
- List of attributes
 - A tuple t of $r(R)$ is denoted by
 $t = (v_1, v_2, \dots, v_n), v_i \in D_i, 1 \leq i \leq n$

SQL Insert

STUDENT(SID, Name, Major, GPA)

- Implicit (list):**
INSERT INTO STUDENT
VALUES (165, 'Susan Jones', 'CS', 0.00);
- Explicit (set):**
INSERT INTO STUDENT (SID, Name)
VALUES (165, 'Susan Jones');

INSERT INTO STUDENT (Name, SID)
VALUES ('Susan Jones', 165);
- Values-clause may be a list of tuples in some systems

Properties of Relations

- A relation is finite
- There are no duplicate tuples in a relation
 - Recall a relation is a set of tuples
- Order of tuples in a relation is not important
 - Many logical orders can be specified on a relation
- A value may appear multiple times in a column
- Order of attribute values in a tuple is
 - important in a *list-of-attributes* definition
 - not important in a *set-of-attributes* definition

Relation Schema

STUDENT

<i>SID</i>	<i>LName</i>	<i>Name</i>	<i>Class</i>	<i>Major</i>
123	Smith	John	3	CS
395	Aiken	Mary	4	CS

← Schema

◆ What is the meaning?

- A **relation schema** R specifies
 - The name of the relation
 - the attribute names A_i of R
 - the domain D_i (data type + format) for each attribute A_i
- data type is a set of **atomic data** values:
 - no attribute is a set-valued (**1st Normal Form**, 1-NF)
 - no attribute is composite
- format specifies the representation of a data value

Example Table Schema

Schema of STUDENT(SID, Name, Major, GPA)

CREATE TABLE STUDENT

```
(  SID  INTEGER,
   Name CHAR(20),
   Major CHAR(4),
   GPA  DEC(3,2)
);
```

Creating a Schema



- Corresponding database is at an **empty** state!
- **Initial state** when the database is **populated** (loaded)
- Domain (type) of each field is **specified** and **enforced** by the DBMS whenever tuples are added or modified

Example: Domain Constraints

<i>SID</i>	<i>Name</i>	<i>Login</i>	<i>Age</i>	<i>GPA</i>
546007	Jones	jones@cs	18	3.4
546100	Smith	smith@ee	18	3.2
546500	Smith	smith@math	19	3.8

□ Example of **IC Violation**:

UPDATE Students

SET Age = `Eighteen`

✗ ✗ ✗

WHERE Name = Jones;

Useful Terms

- **Cardinality** of a relation $r(R)$: # of tuples in $r(R)$ (denoted by $|r(R)|$)
- **Arity** or **degree** of $r(R)$: # of attributes in R (denoted by $|R|$)

$$|R| = 4$$

$$|r(R)| = 3$$

SID	Degree	Major	Year
123	BS	Math	1992
064	BA	History	1991
445	PhD	CS	1999

- ◆ $|R| > 0$ and $|r(R)| \geq 0$
- ◆ Cardinality is property of a relation
- ◆ Arity is property of relation schema or a relation

Relational Database Schema


- A *database schema* is a set of relation schemas and a set of **integrity constraints**



- Integrity Constraints
 - **Structural** Integrity Constraints
 - **key** constraints: uniqueness of keys
 - **entity integrity** constraint: no primary key value can be **NULL**
 - **referential integrity** constraint
 - **Semantic** Integrity Constraints
 - E.g., ??

Integrity Constraints (ICs)

- **IC**: condition that must be true for *any* instance of the database (e.g., domain constraints)
 - A **legal** instance of a relation is one that satisfies all specified ICs
 - ICs are specified when schema is **defined**
 - ICs are enforced when tables are **modified**



Primary Key Constraint

- A set of fields is a **key** for a relation if :
 - No two distinct tuples can have same values in all key fields



- If there is more than one key for a relation:
 - Each is called a candidate key
 - One candidate key is designated as the **primary key**
 - Other candidate key(s) are designated as **alternative** or **unique key(s)**



Example of Keys

<i>SID</i>	<i>Name</i>	<i>Login</i>	<i>Age</i>	<i>GPA</i>
546007	Jones	jones@cs	18	3.4
546100	Smith	smith@ee	18	3.2
546500	Smith	smith@math	19	3.8

- ❑ **Candidate Keys:** *SID*, and *Login*
- ❑ **Primary Key:** *SID*
- ❑ **Unique Key:** *Login*

Example Table Schema in SQL

Schema of STUDENT(*SID*, *Login*, Name, Major, GPA)

```
CREATE TABLE STUDENT
(
  SID INTEGER NOT NULL,
  Login CHAR(15),
  Name CHAR(20),
  Major CHAR(4),
  GPA DEC(3,2),

  CONSTRAINT STUDENT_PK
    PRIMARY KEY (SID),
  CONSTRAINT STUDENT_UN
    UNIQUE (Login)           -- UNIQUE can take NULL values
);
```

Example Table Schema in SQL (2)

Schema of STUDENT(*SID*, Login, Name, *SSN*, GPA)

```
CREATE TABLE STUDENT
(
  SID INTEGER NOT NULL,
  Login CHAR(15),
  Name CHAR(20),
  SSN CHAR(9),
  GPA DEC(3,2),

  CONSTRAINT STUDENT_PK PRIMARY KEY (SID),
  CONSTRAINT STUDENT_UN_SSN
    UNIQUE (SSN),
  CONSTRAINT STUDENT_UN_Login
    UNIQUE (Login)
);
```

Example Table Schema in SQL (3)

Schema of STUDENT(*SID*, *SSN*, Name, *Login*, GPA)

```
CREATE TABLE STUDENT
(
  SID INTEGER
    CONSTRAINT STUDENT_PK_NOT_NULL NOT NULL
    CONSTRAINT STUDENT_PK PRIMARY KEY,
  SSN CHAR(9)
    CONSTRAINT STUDENT_UN_SSN UNIQUE,
  Name CHAR(20),
  Login CHAR(15),
  GPA DEC(3,2),

  CONSTRAINT STUDENT_UN_Login
    UNIQUE (Login)
);
```

Identifying the Key

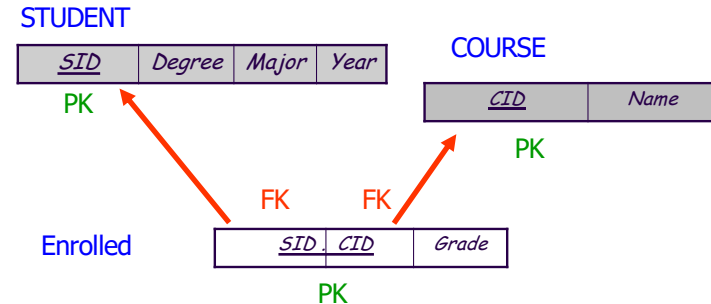
- What is the key in relation GRADUATE=(SID, Degree, Major, Year) ?

<i>SID</i>	<i>Degree</i>	<i>Major</i>	<i>Year</i>
123	BS	CS	1992
123	MS	CS	1993
064	BA	History	1991
445	PhD	CS	1999
123	BS	Math	1992
123	MS	Math	1992




Foreign Keys

- Foreign key (FK)** in relation R_2 is a set of attributes of R_2 that forms a primary key (PK) of another relation R_1 .
 - Attributes in FK and PK have the **same domain**



Foreign Key & Primary Key



- Foreign key:** Set of fields in one relation that is used to “refer” to a tuple in another relation
 - Must correspond to primary key of the referred relation
 - E.g. *SID* is a foreign key referring to *Students*

Foreign Key Constraints

- If foreign key constraints are enforced, **referential integrity** is achieved
 - E.g.: Only students can enroll in a class
 - Only students listed in the “Students” relation should be allowed to enroll for courses
- Like a “*logical pointer*”
 - There shouldn’t be dangling references
 - Either valid PK or NULL

Any Attribute can be a Foreign Key

Faculty			Courses		
<u>FID</u>	Name	Area	<u>CID</u>	Name	Instructor
007	Panos	DB	CS1555	DB	007
100	Daniel	OS	CS1530	SW	NULL
500	Adriana	AI	CS1550	OS	100


Primary Key

Primary Key

Foreign Key

- Foreign key: Set of fields in one relation that is used to “refer” to a tuple in another relation
 - Must correspond to primary key of the referred relation
 - If not part of a key, it could be NULL

Foreign Keys in SQL



```

CREATE TABLE Enrolled (
    SID CHAR(20), CID CHAR(20), Grade CHAR(2),
    CONSTRAINT Enrolled_PK PRIMARY KEY (SID, CID),
    CONSTRAINT Enrolled_FK_sid
        FOREIGN KEY (SID) REFERENCES Students (SID),
    CONSTRAINT Enrolled_FK_cid
        FOREIGN KEY (CID) REFERENCES Courses
);
    
```

Referential Integrity Constraints

<u>SID</u>	Name	Age	<u>SID</u>	<u>CID</u>	Grade	<u>CID</u>	CName
546007	Juan	18	546007	CS1555	A	CS1555	DB
546100	Bob	19	546007	CS1530	B+	CS1530	SW
546500	Bill	20	546500	CS1555	B	CS1550	OS

Students

Enrolled

Courses

- Any IC Violation?
- ```

DELETE
FROM Enrolled
WHERE SID = 546500;

No Violations! ✓

```

## Referential Integrity Constraints

| <u>SID</u> | Name | Age | <u>SID</u> | <u>CID</u> | Grade | <u>CID</u> | CName |
|------------|------|-----|------------|------------|-------|------------|-------|
| 546007     | Juan | 18  | 546007     | CS 1555    | A     | CS 1555    | DB    |
| 546100     | Bob  | 19  | 546007     | CS1530     | B+    | CS1530     | SW    |
| 546500     | Bill | 20  | 546500     | CS 1555    | B     | CS 1550    | OS    |

Students

Enrolled

Courses

- Any IC Violation?
- ```

DELETE
FROM Students
WHERE SID = 546500;

Violation! ✗
    
```



Referential Integrity Enforcement

- ❑ What are the alternatives when a “Students” tuple is **deleted**?
 1. **Delete all** Enrolled tuples that refer to it
 2. **Disallow** deletion of a Students tuple that is referred to
 3. **Set SID** in Enrolled tuples that refer to it to some “default” SID (e.g., 000000)
 4. If SID was not part of the primary key, **Set SID** to a special value “NULL”, denoting “unknown” or “inapplicable”

Referential Integrity in SQL

- ❑ SQL/92 and SQL/99 support all 4 options on delete and update:
 - **NO ACTION** (default)
 - delete/update is rejected
 - **CASCADE**
 - also delete all tuples that refer to deleted tuple
 - **SET NULL / SET DEFAULT**
 - sets foreign key value of referencing tuple

RI Trigger Actions in SQL



- ❑ **CREATE TABLE Enrolled** (
 - SID CHAR(20), CID CHAR(20), Grade CHAR(2),
 - CONSTRAINT Enrolled_PK PRIMARY KEY (SID, CID),
 - CONSTRAINT Enrolled_FK_sid
 - FOREIGN KEY (SID) REFERENCES Students (SID),
 - CONSTRAINT Enrolled_FK_cid
 - FOREIGN KEY (CID) REFERENCES Courses
 - ON UPDATE CASCADE ON DELETE NO ACTION

Enforcing Integrity Constraints



- ❑ What would be the outcome?
 - Insert (585811, 'Jie', 19, 3.95) into Students
 - Insert (585811, NULL, NULL) into Enrollment
 - Insert (546100, 'CS 1555', NULL) into Enrollment
 - Insert (546100, 'Mary', 18, 3.65) into Students
 - Delete ('CS 1530') from Courses

SID	Name	Age	GPA
546007	Susan	18	3.8
546100	Bob	19	3.65
546500	Bill	20	3.7

Students

CID	Name
CS1555	DB
CS1530	SW
CS1550	OS

Courses

SID	CID	Grade
546007	CS1550	A
546007	CS1530	B+
546100	CS1550	B

Enrollment