

Scalar Q2 Execution

Students

SID	Name	Class	Major
123	John	3	CS
124	Mary	3	CS
126	Sam	2	CS
999	Newman	1	CS
132	Erin	2	EE

Enroll

SID	CID	Term	Grade
123	CS1520	Fall 10	3.75
124	CS1520	Fall 10	4
126	CS1520	Fall 10	3
123	CS1555	Fall 10	4
124	CS1555	Fall 10	NULL
126	CS1550	Spring 11	NULL

Q2 RESULT

SID	LG
123	4
124	4
126	3
999	

- Q2 is equivalent to a single-level SQL query; which one?

Equivalent Query to Q2

Q2 Equivalent Query:

```
SELECT SID, MAX(Grade) AS LG
FROM STUDENT S NATURAL LEFT JOIN ENROLL E
WHERE S.Major = 'CS'
GROUP BY SID;
```

Set Membership $x \in A$

- The IN and NOT IN operators check for simple membership.
 - LIBRARIAN (SSN, Name, City, Gender, Salary, SNO);
SECTION (SNO, Name, HeadSSN);
- Q1: List each head librarian's SSN, along with their section, except those of the Science & Art sections.
- ```
SELECT HeadSSN, Name
FROM SECTION
WHERE Name NOT IN ('Science', 'Art');
```

## Set Membership

LIBRARIAN (SSN, Name, City, Gender, Salary, SNO);  
SECTION (SNO, Name, HeadSSN);  
DEPENDENT (Name, LIBSSN, DSSN, Gender, DoB);

- Q2: ? Lists librarian's name who have dependents of the same gender

```
SELECT L.Name
FROM LIBRARIAN AS L
WHERE L.SSN IN (SELECT LIBSSN
FROM DEPENDENT D
WHERE D.LIBSSN = L.SSN AND
D.Gender = L.Gender);
```

## Set Membership & Comparisons

- ❑ Test for membership on other comparisons:  
( =, <>, >, >=, <, <= )
- ❑ They can be quantified using **ANY** (i.e., SOME) or **ALL**.
- ❑ Q: List the SSN of all head librarians whose salary is lower than that of any librarian who is not a head librarian.

```
SELECT H.SSN
FROM LIBRARIAN H JOIN SECTION ON SSN = HeadSSN
WHERE H.Salary < ANY (SELECT L.Salary
 FROM LIBRARIAN L
 WHERE L.SSN NOT IN (SELECT HeadSSN
 FROM SECTION));
```

## Set Comparisons: Unique & Empty

- ❑ The UNIQUE and NOT UNIQUE operators test for duplicates in a result (i.e, tests for a set and a bag, respectively)
- ❑ The EXISTS (not empty) and NOT EXISTS (empty) operators test for emptiness of a result

- ❑ Q: ? Lists non-head librarian

```
SELECT L.SSN, L.Name
FROM LIBRARIAN L
WHERE NOT EXISTS (SELECT *
 FROM SECTION
 WHERE L.SSN = HeadSSN);
```

## Set Comparisons: Unique & Empty...

- ❑ Q: ? Lists the students with dual major CS and Math

```
SELECT S.SID
FROM STUDENT S
WHERE NOT UNIQUE (SELECT *
 FROM (SELECT SID
 FROM STUDENT
 WHERE Major = 'CS') UNION ALL
 (SELECT SID
 FROM STUDENT
 WHERE Major = 'Math')
 WHERE S.SID = SID
);
```

## Unique & Not Unique

- ❑ UNIQUE in POSTGRES & Oracle is used for constraints or instead of DISTINCT
  - Always use **DISTINCT**
- ❑ UNIQUE & Not UNIQUE can be expressed using JOIN

- ❑ Q: ? Lists students with dual major only once

```
SELECT C.SID
FROM Student C JOIN student M ON
 (C.SID = M.SID AND C.Major > M.Major);
```

## Unique in Postgres/Oracle

```
SELECT S.SID
FROM STUDENT S
WHERE NOT EXISTS (SELECT S.SID
 FROM ((SELECT SID
 FROM STUDENT
 WHERE Major = 'CS')
 UNION ALL
 (SELECT SID FROM STUDENT WHERE
 Major = 'Math')) A
 WHERE S.SID = SID
 GROUP BY S.SID
 HAVING count(*) > 1
);
```

## Removing Duplicates

- ❑ Q: List all the students who are double majors with their majors only once.
- ❑ Note that (1, CS, Math) and (1, Math, CS) are duplicates
- ❑ A: **SELECT C.SID, C.Major, M.Major**  
**FROM Student C JOIN student M ON**  
**(C.SID = M.SID AND C.Major > M.Major);**

❑ Example:

**Student**

| SID | Major |
|-----|-------|
| 123 | CS    |
| 123 | Math  |

**Student C JOIN Student M**

| C.SID | C.Major | M.SID | M. Major |
|-------|---------|-------|----------|
| 123   | CS      | 123   | CS       |
| 123   | CS      | 123   | Math     |
| 123   | Math    | 123   | CS       |
| 123   | Math    | 123   | Math     |

Result in bold

## Challenging yet Common Query

- ❑ Assume ENROLL(SID, CID, score)
- ❑ Find the ranking of students in CS2550 according to their scores. Your results should consider the case of tie.
  - E.g., output when 007 & 009 both received the highest score.

| SID | Rank |
|-----|------|
| 007 | 1    |
| 009 | 1    |
| 003 | 3    |
| 005 | 4    |

## Ranking Query using of Scalar Subquery

- ❑ ENROLL(SID, CID, score)
- ❑ Q: Find the ranking of students in CS1555 according to their scores. Your results should consider the case of tie.

```
SELECT S.SID, (1 + (SELECT COUNT(*)
 FROM ENROLL E
 WHERE E.CID = 'CS 1555' AND
 E.score > S.score)
) AS Rank
FROM ENROLL S
WHERE S.CID = 'CS 1555'
ORDER BY Rank;
```

## Limiting result rows

- ❑ Save resources, speed-up result
- ❑ FETCH clause (SQL 2008)
  - OFFSET *start* { ROW | ROWS }
  - FETCH { FIRST | NEXT } [ *count* ] { ROW | ROWS } ONLY
  - the **OFFSET** clause must come before the **FETCH** clause
  - *start*: number of rows to skip [default 0]
  - *count*: maximum number of rows to return [default 1]
  - *ROW*, *ROWS* and *NEXT* is “noise” keywords – have no effect
- ❑ E.g.,

```
SELECT *
FROM STUDENT
FETCH FIRST 10 ROWS ONLY;
```

## Top-K Queries

- ❑ Q1: ? Ten “oldest” students

```
SELECT *
FROM STUDENT
ORDER BY SID ASC
FETCH FIRST 10 ROWS ONLY;
```
- ❑ Q2: ? Ten students with lowest QPA

```
SELECT *
FROM STUDENT
ORDER BY QPA ASC
FETCH FIRST 10 ROWS ONLY;
```
- ❑ How many tuples does Q2 return?

## Next-K Queries

- ❑ Q3: ? List students with lowest QPA rank between 7-16

```
SELECT *
FROM STUDENT
ORDER BY QPA ASC
OFFSET 6 ROW
FETCH NEXT 10 ROWS ONLY;
```
- ❑ Q3: (alternative that seems to work)

```
SELECT *
FROM STUDENT
ORDER BY QPA ASC
OFFSET 6 ROW
FETCH FIRST 10 ROWS ONLY;
```

## Limits (No Standard Syntax)

- ❑ `SELECT * FROM T WHERE ROWNUM <= 10;`
  - Oracle (also supports the standard ?)
- ❑ `SELECT * FROM T LIMIT 10 OFFSET 20;`
  - MySQL, PostgreSQL (supports the standard), SQLite
- ❑ `SELECT * FROM T WHERE ID_T > 10 FETCH FIRST 10 ROWS ONLY;`
  - IBM DB2
- ❑ `SELECT TOP 10 * FROM T;`
  - MS SQL Server, Sybase ASE

## Top-K in Oracle

- Q: Select the top 3 students with the highest QPA

```
SELECT *
FROM (SELECT * FROM STUDENT ORDER BY QPA DESC) S
WHERE rownum <= 3
ORDER BY rownum;
```

- ROWNUM: pseudocolumn which returns a number showing the order in which Oracle selects a row from a table
- Q: Select top students with offset = 3, next = 3  
??

## Top-K in Oracle

- Q: Select top students with offset = 3, next = 3
- What about using BETWEEN?

```
SELECT *
FROM (SELECT * FROM STUDENT ORDER BY QPA DESC)
WHERE rownum BETWEEN 3 and 6;
```

Or

```
SELECT *
FROM (SELECT * FROM STUDENT ORDER BY QPA DESC)
WHERE rownum > 3 AND rownum <= 6;
```

- Does not work because rownum is assigned dynamically

## Next Top-K in Oracle

- Q: Select top students with offset = 3, next = 3

```
SELECT *
FROM (
 SELECT SID, Name, Major, QPA, rownum AS snum
 FROM (SELECT * FROM STUDENT ORDER BY QPA DESC)
 WHERE rownum <= 6)
WHERE snum > 3;
```

- Consider QPAs: 1, 3, 6, 8, 2, 12, 19, 9

## Top-K in Postgres

```
SELECT *
FROM (
 SELECT *,
 ROW_NUMBER() OVER (ORDER BY sid)
 FROM cs1555.student
) x
WHERE ROW_NUMBER <= 3;
```

## Structured Query Language SQL – DML

### ◆ Update Operations

## Query Language with Update Statements?

- Is it a historical accident?  
or Select is in the hard of updates!
- *Selections* are expressed  
by the WHERE clause!



## Update Tuples

- Update can apply to a **single** relation
- Updates all the selected tuples by the condition in the WHERE-clause

- Examples:

```
UPDATE STUDENT
SET Name = 'Kathy Jones'
WHERE SID = 165;
```

```
UPDATE STUDENT
SET Major = 'CS'
WHERE DNO IN
(SELECT DNUM
 FROM DEPT
 WHERE Dname = 'CS');
```

## Delete Tuple

- Delete removes all selected tuples by the condition in the WHERE-clause
- Examples:

```
DELETE FROM STUDENT
WHERE SID = 165;
```

```
DELETE FROM STUDENT
WHERE Name = 'John';
```

```
DELETE FROM STUDENT;
```

```
DELETE FROM STUDENT
WHERE DNO IN
(SELECT DNUM
 FROM DEPT
 WHERE Dname = 'CS');
```

## Insert

`STUDENT(SID,Name,Major);`

- ❑ Two forms: **Implicit** (list) and **Explicit** (set)

- ❑ **Implicit:**

```
INSERT INTO STUDENT
VALUES (165, 'Susan', 'CS');
```

- ❑ **Explicit:**

```
INSERT INTO STUDENT (SID, Name)
VALUES (165, 'Susan Jones');

INSERT INTO STUDENT (Name, SID)
VALUES ('Susan Jones', 165);
```

## Derived Insert Values

- ❑ Tuples are derived using SELECT
- ❑ Useful to populate a table in the database from data already in the database

- ❑ E.g.,

```
INSERT INTO Dept_Info (Dept_Name, Num_Students)
SELECT Major, Count(*)
FROM STUDENT
GROUP BY Major;
```