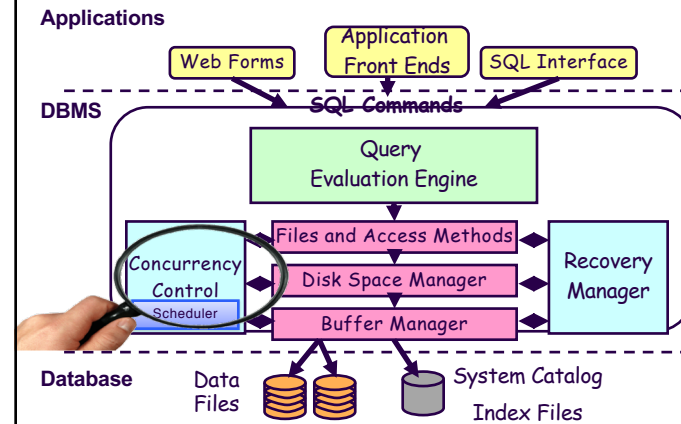


Transaction Processing: Concurrency control

Structure of a DBMS




ACID Properties

Atomicity
Consistency
Isolation
Durability



ACID Properties

Property	Dealt with by
A, D	Recovery Techniques
I	Concurrency Control Techniques
C	Checks, Assertions, Triggers Applications Programmers



Two Views of the System

```
set transaction read write
select * from Students
insert into Students values (777, 'Jane', 'CS')
Commit;
```

1. Application Programmer's View
 - Start
 - sequence of **SQL statements**
 - Commit or Rollback
2. System developer's View
 - Start
 - sequence of **Reads and Writes**
 - Commit or Abort

Transactions

T₁: UPDATE Accounts SET balance= balance + 100
WHERE client=7

T₂: UPDATE Accounts SET balance= balance + 500
WHERE client=7

- Assume that initially, balance = \$1000
- What is the balance after executing T₁ & T₂?
 - should be \$1600

However things might go wrong!!

Interleaved Transactions

T₁: UPDATE Accounts SET balance= balance + 100 WHERE client=7

T₂: UPDATE Accounts SET balance= balance + 500 WHERE client=7

- Update (balance) =
 - Read (balance); Modify (balance); Write (balance)
- Again, assume that initially balance = \$1000
- What happens if T₁ and T₂ are executed **concurrently** and they both issue Read (balance) at the same time?
 - If T₁ finishes last, balance = \$1100
 - If T₂ finishes last, balance = \$1500
 - And both values are **incorrect!**



Isolation

T₁: UPDATE Accounts SET balance= balance + 100 WHERE client=7

T₂: UPDATE Accounts SET balance= balance + 500 WHERE client=7

- **Isolation:** The result of the execution of **concurrent** transactions is the same as if transactions were executed **serially** (one after the other)
- **Serializability:** Operations may be interleaved, but execution must be equivalent to some sequential (**serial**) order of transactions
 - E.g., T₁ followed by T₂, or T₂ followed by T₁
- Mechanism: **Concurrency Control**



Serial Executions

Item X = 100

T ₁	T ₂	T ₁	T ₂
a=read_item(X) a=a+10 write_item(X, a) commit			b=read_item(X) b=b*2 write_item(X, b) commit
	b=read_item(X) b=b*2 write_item(X, b) commit	a=read_item(X) a=a+10 write_item(X, a) commit	

Item X = 220

Item X = 210

Concurrency Goal

- ❑ **Concurrency Goal:** Execute a sequence of SQL statements so they “appear” to be running in **isolation**

- ❑ Simple Solution

- **Execute** them in isolation!

- ❑ But want to enable concurrency whenever it is **safe**:

- High performance DBMS
- Benefit from modern architectures (e.g., multicore processors, etc.)



Anomalies

- ❑ **Question:** why **concurrency control** is needed?

- ❑ **Answer:** to avoid the following anomalies:

1. The **lost update** problem
2. The **dirty read** problem
3. The **unrepeatable read** problem
 - The **phantom read**



Three Bad Dependencies

- ❑ **Lost Update:** Read_i(X) Write_j(X) Write_i(X) sequence
 - Write-Write interaction (W-W)
- ❑ **Dirty Data:** Write_i(X) Read_j(X) sequence
 - Write-Read interaction (W-R)
- ❑ **Unrepeatable Read:** Read_i(X) Write_j(X) Read_i(X) sequence
 - Read-Write interaction (R-W)
- ❑ These forms of inconsistency are the whole story.

Conflicting Operations

- ❑ A conflict happens if we have two operations such that:
 1. they belong to two different transactions, and
 2. they both operate on the **same data item**,
 3. and **one of them is a write**



Conflicting Operations

- ❑ Two operations **conflict** if it matters in which order they are performed
 - The order affects the results;
 - The order affects the state of the database.
- ❑ Non conflicting operations are called **compatible**.
- ❑ A **compatibility table** shows which operations are compatible.
- ❑ E.g., {Read, Write}

	Read	Write
Read	yes	no
Write	No	no

Schedules

- ❑ When transactions are executing concurrently, the order of execution of operations from all transactions is known as a **schedule** (or **history**)
- ❑ A Schedule **S** of **n** transactions T_1, T_2, \dots, T_n is an ordering of the operations of the transactions
- ❑ For the purpose of concurrency control, we are mainly interested in the **read (r)** and **write (w)** operations, as well as **commit (c)** and **abort (a)** operations

Schedule: example 1

T_1	T_2	Timeline
read_item(X) $X = X - N$	read_item(X) $X = X + M$	<div>Mapping:</div> <ul style="list-style-type: none"> • Drop local variables, e.g., a, b, c... • Use db items names, e.g., X and Y, to replace local variables • Project on the time line
write_item(X) read_item(Y)	write_item(X) commit	
$Y = Y + N$ write_item(Y) commit		

S: $r_1(X), r_2(X), w_1(X), r_1(Y), w_2(X), c_2, w_1(Y), c_1$

Scheduling Transactions

- ❑ **Serial schedule:** Schedule that does not interleave the actions of different transactions
- ❑ **Serializable schedule:** A schedule that is equivalent to some serial execution of the transactions
- ❑ **Result Equivalent schedules:** For any database state, two schedules S_1 and S_2 are equivalent if:
 - the state produced by executing S_1 is identical to the state produced by executing S_2

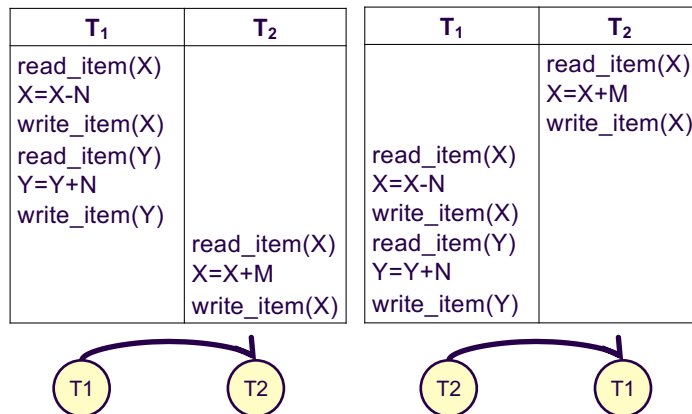
How to achieve serializability? Concurrency Control

- **Equivalence** is defined wrt conflicting operations: the order of any two conflicting operations is the same in S_1 and S_2

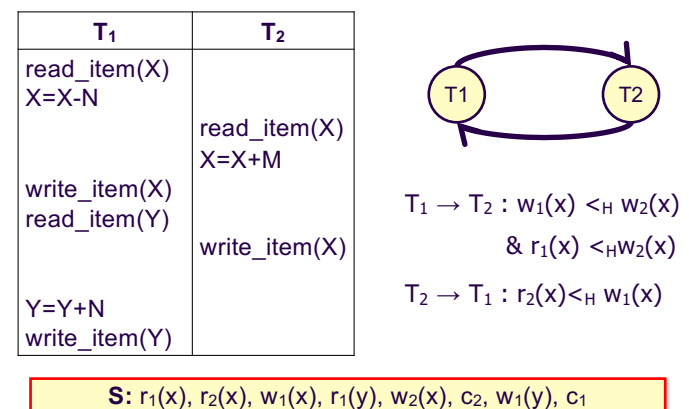
Testing CSR

- ❑ Test for CSR by analyzing the precedence graph **SG(H)**, called a **serialization graph**, derived from history H.
- ❑ An SG(H) is a directed graph in which:
 1. nodes represent transactions in H;
 2. an edge $T_i \rightarrow T_j$, $i \neq j$, means that one of T_i 's operations precedes and conflicts with one of T_j 's operations in H.
- ❑ **Serializability Theorem:**
A history H is serializable iff SG(H) is **acyclic**.
- ❑ Proof: ?

Testing CSR: Example



Testing CSR: Example



Testing CSR: Example

T ₁	T ₂
read_item(X) X=X-N write_item(X)	
	read_item(X) X=X+M write_item(X)
read_item(Y) Y=Y+N write_item(Y)	



S: $r_1(x), w_1(x), r_2(x), w_2(x), c_2, r_1(y), w_1(y), c_1$

Finding The Equivalent Serial History

- Make a topological sorting of the serialization graph.

$H = r_1(x) w_1(x) r_2(y) w_2(y) c_2 r_1(y) w_1(y) c_1 r_3(x) w_3(x) c_3$

SG (H) 

$H_s = r_2(y) w_2(y) c_2 r_1(x) w_1(x) r_1(y) w_1(y) c_1 r_3(x) w_3(x) c_3$

T₂

T₁


T₃

Finding The Equivalent Serial History ...

- Several serial orders can be obtained by topological sorting:

$H' = r_1(x) w_1(x) r_3(x) w_2(y) c_1 r_3(y) c_2 w_3(y) c_3$

SG (H')



$H_s: T_1 T_2 T_3 \text{ or } T_2 T_1 T_3$

Three Bad Dependencies

- Lost Update:** $\text{Read}_i(X) \text{ Write}_j(X) \text{ Write}_i(X)$ sequence
 - Write-Write interaction (W-W)
- Dirty Data:** $\text{Write}_i(X) \text{ Read}_j(X)$ sequence
 - Write-Read interaction (W-R)
- Unrepeatable Read:** $\text{Read}_i(X) \text{ Write}_j(X) \text{ Read}_i(X)$ sequence
 - Read-Write interaction (R-W)
- These forms of inconsistency are the whole story.

ANSI SQL2 Isolation Levels

- ❑ SET TRANSACTION READ ONLY | READ WRITE
[ISOLATION LEVEL READ UNCOMMITTED |
READ COMMIT |
REPEATABLE READ |
SERIALIZABLE]

Concurrency Control Schemes

- ❑ Lock-based CC schemes
 - **Two-phase locking** [IBM DB2, SQLServer]
 - Multigranularity locking
 - Tree/Index locking
- ❑ **Multiversion** [Oracle, SQLServer]
- ❑ Timestamp-based
- ❑ Optimistic CC & Certifiers

Lock Based Concurrency Control

- ❑ Locking is the most common **synchronization** mechanism
- ❑ A **lock** is associated with each data item in the database
- ❑ A lock on item “**x**” indicates that a transaction is **performing** an operation (read or write) on “**x**”.



Lock Based Concurrency Control

- ❑ Transaction T_i can issue the following operations on item x :

- **read_lock (x)**

- x is read-locked by T_i
- **shared** lock: other transactions are allowed to read x



- **write_lock (x)**

- x is write-locked by T_i
- **exclusive** lock: single transaction holds the lock on x



- **unlock (x)**

Basic Two Phase Locking (2PL)

- A scheduler following the 2PL protocol has two phases:

1. A Growing phase

- Whenever the scheduler receives an operation on any item, it must **acquire** a lock on that item before executing the operation.
- No locks can be released in this phase

2. A Shrinking phase

- Once a scheduler has **released** a lock for a transaction, it cannot request any additional locks on any data item for this transaction.

Basic Two Phase Locking (2PL)

- Example:

- **Transaction T:** $a = r(x); w(y, a);$

S₁: read_lock(x); a=r(x); write_lock(y); w(y, a); unlock(x); unlock(y); ✓

S₂: read_lock(x); a=r(x); unlock(x); write_lock(y); w(y, a); unlock(y); ✗

S₃: read_lock(x); a=r(x); write_lock(y); unlock(x); w(y, a); unlock(y); ✓

Rigorous 2PL or industrial Strict 2PL


- The growing phase
 - transactions request locks just before they operate on a data item.
- The growing phase ends at *commit time*.
 - no locks can be released until commit or abort time.
 - no overwriting of dirty data.
 - no overwriting of data read by active transactions.
 - no reading of dirty data.
- Easy to implement a strict 2PL. Why ?
- Has a functional advantage. What ?



Issues Related to Locking



Deadlock



Livelock



Starvation

Deadlocks

Examples:

(I) 2 Items			(II) 1 Item		
T_1	T_2	Comments	T_1	T_2	Comments
rl(x)		granted	rl(x)		granted
	rl(y)	granted		rl(x)	granted
wl(y)		T_1 blocked	wl(x)		T_1 blocked
	wl(x)	T_2 blocked (deadlock)		wl(x)	T_2 blocked (deadlock)

- Example II involves lock conversion
- The scheduler *restarts* any transaction aborted due to deadlock.

Deadlocks

- A *deadlock* occurs when two or more transactions are blocked indefinitely.
- This happens because each holds locks on data items on which the other transaction(s) attempt to place a conflicting lock.
- Necessary conditions for deadlock situations.
 - mutual exclusion
 - hold and wait
 - no preemption
 - circular wait.


Deadlock Handling Schemes

- Deadlock avoidance
 - Timestamp ordering (Wait-Die, Wound-Wait)
- Deadlock Prevention
 - Predeclaration of resources,...
- Deadlock Detection and Resolution
 - Time-out
 - Wait-for graphs

Deadlock Detection

Deadlock Detection:

- Deadlocks are **allowed** to happen!
- A wait-for-graph is used for detecting **cycle**



wait-for-graph (WFG)

- Created using the **lock table**
- As soon as a transaction is blocked, it is added to the graph
- Detect cycles:** T_i waits for T_j and T_j waits for T_i , then this creates a cycle!
- One of the transactions in a cycle is **rolled back (aborted)**
 - Which one?

Jim Gray - the Godfather of Transactions!



Concurrency Control Schemes

- Lock-based CC schemes
 - Two-phase locking** [IBM DB2, SQLServer]
 - Multigranularity locking
 - Tree/Index locking
- Multiversion** [Oracle, SQLServer]
- Timestamp-based
- Optimistic CC & Certifiers

Multiversion Concurrency Control

- Assume the following sequence of events.
 $W_0(x) C_0 W_2(x) R_1(x) C_2 C_1$
- This sequence CANNOT be produced by a strict 2PL scheduler
 - T_1 can not read lock x until after C_2
- An Idea !!
 - If we had kept the old version of x when $W_2(x)$, then we could avoid having to delay T_1 as in 2PL by having T_1 read the previous (old) value of x (produced by T_0).

Basic Idea

- ❑ The DBMS keeps a list of versions for each x
 - Version x_i means the version of x produced by a Write on x by transaction T_i
- ❑ Each Write(x) produces a new version of x
- ❑ When the scheduler receives a $R_i(x)$, it must decide which version of x to read
 - A Read operation will be converted to the form $R(x_i)$
- ❑ If a transaction T is aborted, any version it created is destroyed
- ❑ If a transaction T is committed, any version it created becomes available for reading by other transactions

CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

75

Two Version 2PL (2V2PL)

- ❑ keep one or two versions of each data item x .
- ❑ When a T_i wants to write x , it sets a $wl(x)$ and creates a new version of x , x_i .
 - The $wl(x)$ prohibits other transactions from writing x .
- ❑ Readers are allowed to place a $r/$ on their write-locked x or the previous version of x .
- ❑ When T_i commits, the x_i version of x becomes x 's unique version (the previous x may now be deleted).
- ❑ To delete the previous x when T_i commits, we need to know that no other transaction reads x .
 - Request a commit lock (cl) which conflict with rl
- ❑ Deadlocks are possible and indicate non-CSR execution
 - use any deadlock detection or prevention technique.

CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

76

Example Transaction

- ❑ **Class** (classid, max_num_students, cur_num_students)
- ❑ Consider transaction "Enroll_student"

```
SET TRANSACTIONS READ WRITE;  
SELECT max_num_students, cur_num_students  
FROM CLASS  
WHERE classID = 1;
```

Read(classid = 1)

```
If (cur_num_students < max_num_students)  
  update CLASS  
  set cur_num_students = cur_num_students + 1  
  where classID = 1;  
else  
  print 'the class is full';  
COMMIT;
```

Write(classid = 1)

CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

77

Concurrent Transactions

```
SET TRANSACTIONS READ;  
SELECT max_num_students, cur_num_students  
FROM CLASS  
WHERE classID = 1;  
  
sleep...  
  
If (cur_num_students < max_num_students)  
  update CLASS  
  set cur_num_students = cur_num_students + 1  
  where classID = 1;  
else  
  print 'the class is full';  
COMMIT;
```

CS1555/2055, Panos K. Chrysanthis – University of Pittsburgh

78

Concurrent Transactions R/Ws

- Assume `max_num_students = 40`, `cur_num_students = 39`
- Execution:
 - `r1(max_num_students)`
 - `r1(cur_num_students)` -- `cur_num_students = 39`
 - ... `sleep1`
 - `r2(max_num_students)`
 - `r2(cur_num_students).` -- `cur_num_students = 39`
 - ... `sleep2`
 - If (`cur_num_students < max_num_students`)
 - `w1(cur_num_students++)` -- `cur_num_students = 40`
 - If (`cur_num_students < max_num_students`)
 - `w1(cur_num_students++)` -- `cur_num_students = 41`

Write/Exclusive Lock

- Example:
 - `SELECT max_num_students, cur_num_students`
 - `FROM CLASS`
 - `WHERE classID = 1555`
 - `FOR UPDATE OF cur_num_students;`
- Alternative just specify **FOR UPDATE;**
- Error Messages:
 - No lock: Cannot serialize access for this transaction"
 - With lock: "The class is full"

Postgres Isolation Levels

- `SET TRANSACTION READ ONLY | READ WRITE`
[ISOLATION LEVEL ~~READ UNCOMMITTED~~ |
 READ COMMIT |
 REPEATABLE READ |
 SERIALIZABLE]
- `READ COMMITTED` is the *default*
 - Not always the most recent/latest one
 - It cannot see even its own uncommitted updates
- `REPEATABLE READS` always, Why?
- JDBC: `dbcon.TRANSACTION_READ_COMMITTED`,
 `dbcon.TRANSACTION_SERIALIZABLE`

Oracle Isolation Levels

- `SET TRANSACTION READ ONLY | READ WRITE`
[ISOLATION LEVEL **READ COMMITTED** ||
 SERIALIZABLE]
- `READ COMMITTED` is the *default*
 - Not always the most recent/latest one
- `REPEATABLE READS` always, Why?
- JDBC: `dbcon.TRANSACTION_READ_COMMITTED`,
 `dbcon.TRANSACTION_SERIALIZABLE`